

# STAT40800 - Final Project

Suraj Bodhanandan Nhattuvetty - 23200338

```
In [139]: #Import the necessary libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
import numpy as np
```

## Question 1

### a) Load the dataset

```
In [140]: male = pd.read_csv("male_stud.csv")
male
```

Out[140]:

	large_family	lives_in_city	traveltime	studytime	failures	paid
0	0	1	1	2	0	1
1	0	1	1	2	0	0
2	0	1	1	2	0	1
3	1	1	1	2	0	1
4	0	1	1	1	0	1
...	...	...	...	...	...	...
182	0	1	1	2	2	1
183	0	1	2	1	0	0
184	1	0	1	1	3	0
185	0	0	3	1	0	0
186	0	1	1	1	0	0

187 rows × 14 columns



### b) Inspect the data

```
In [141]: #Get the dimensions of the dataframe
male.shape
```

Out[141]: (187, 14)

The dataset includes the data of 187 students. There are 14 different indicators that are used to describe the data.

```
In  list(male.columns)
[142]:
```

```
Out[142]: ['large_family',
 'lives_in_city',
 'traveltime',
 'studytime',
 'failures',
 'paid',
 'activities',
 'internet',
 'romantic',
 'famrel',
 'freetime',
 'goout',
 'absences',
 'final_grade']
```

The 14 different indicators can be seen as above

```
In  #Check if there are any missing values
[143]: male.isnull().values.any()
         male.isna().sum()
```

```
Out[143]: large_family      0
           lives_in_city     0
           traveltime        0
           studytime         0
           failures          0
           paid               0
           activities         0
           internet          0
           romantic          0
           famrel             0
           freetime           0
           goout              0
           absences           0
           final_grade        0
           dtype: int64
```

There are no missing values in the given dataset.

### c) Exploratory data analysis

The indicators in the dataset are in the form numeric categories. A summary of them can be found in the following way

In male.describe()  
[144]:

	large_family	lives_in_city	traveltime	studytime	failure
count	187.000000	187.000000	187.000000	187.000000	187.000000
mean	0.668449	0.764706	1.491979	1.764706	0.368984
std	0.472034	0.425321	0.750405	0.808713	0.788152
min	0.000000	0.000000	1.000000	1.000000	0.000000
25%	0.000000	1.000000	1.000000	1.000000	0.000000
50%	1.000000	1.000000	1.000000	2.000000	0.000000
75%	1.000000	1.000000	2.000000	2.000000	0.000000
max	1.000000	1.000000	4.000000	4.000000	3.000000



It is a good idea to describe the data in the below manner as most of the values act like a categorical data.

In male.astype('object').describe().transpose()  
[145]:

	count	unique	top	freq
large_family	187	2	1	125
lives_in_city	187	2	1	143
traveltime	187	4	1	118
studytime	187	4	2	85
failures	187	4	0	144
paid	187	2	0	114
activities	187	2	1	105
internet	187	2	1	159
romantic	187	2	0	134
famrel	187	5	4	88
freetime	187	5	3	64
goout	187	5	3	60
absences	187	23	0	52
final_grade	187	17	10	26

Looking at the dataset, most male students have a family with more than 3 members. Most also choose to live in urban areas. Most also opt to stay in areas which takes less than 15 minutes to travel to school. On an average they spend 2-5 hours weekly on their studies. Most students do not have any past failures and did not opt for extra paid classes. Most also did take part in extra curricular activities. Almost everyone had internet access at home. Most students were not in a romantic relationship.

Approximately 50 % of the students had a very good relation with their family. On an average, students had a good amount of freetime. Students do go out with their friends for a decent amount of time.

The average number of absences is approximately 5 days with a standard deviation of approximately 6 days. The maximum number of absences is 38 and the minimum is 0. The average grade of student approximately 11 with a standard deviation of approximately 4.5. The minium grade is 0 and maximum is 20.

```
In #Graphical summaries
[146]: fig = plt.figure(figsize=(25,20))
plt.subplot(5,3,1)
sns.countplot(x="large_family", data=male)

plt.subplot(5,3,2)
sns.countplot(x="lives_in_city", data=male)

plt.subplot(5,3,3)
sns.countplot(x="traveltime", data=male)

plt.subplot(5,3,4)
sns.countplot(x="studytme", data=male)

plt.subplot(5,3,5)
sns.countplot(x="failures", data=male)

plt.subplot(5,3,6)
sns.countplot(x="paid", data=male)

plt.subplot(5,3,7)
sns.countplot(x="activities", data=male)

plt.subplot(5,3,8)
sns.countplot(x="internet", data=male)

plt.subplot(5,3,9)
sns.countplot(x="romantic", data=male)

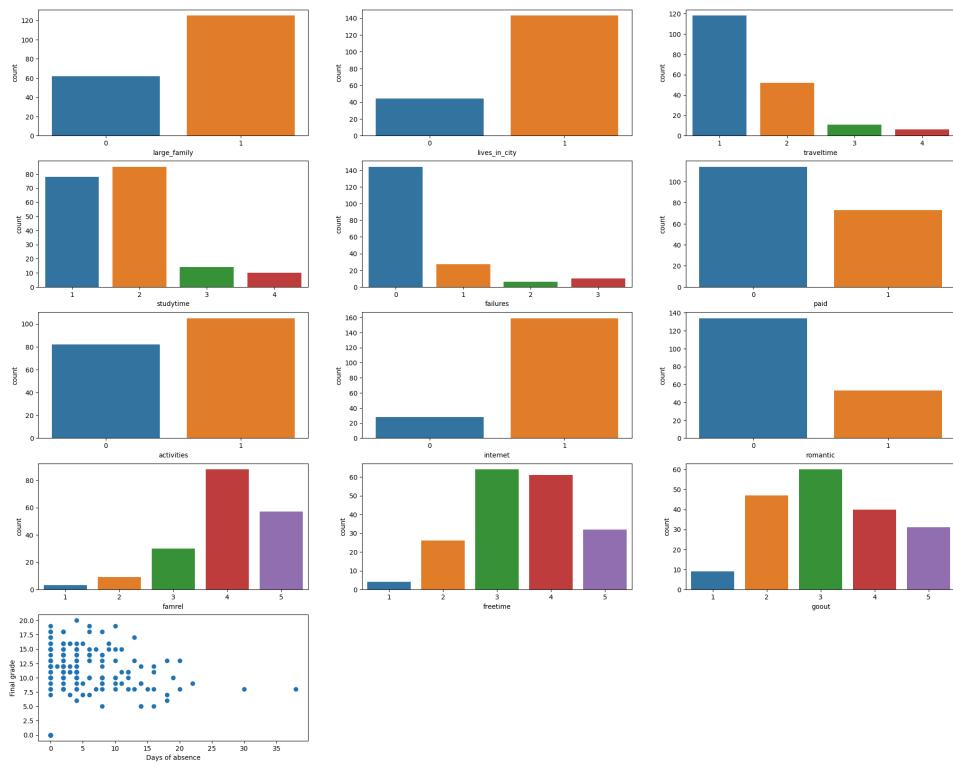
plt.subplot(5,3,10)
sns.countplot(x="famrel", data=male)

plt.subplot(5,3,11)
sns.countplot(x="freetime", data=male)

plt.subplot(5,3,12)
sns.countplot(x="goout", data=male)

plt.subplot(5,3,13)
plt.scatter(male['absences'], male['final_grade'])
plt.xlabel('Days of absence')
plt.ylabel('Final grade')
```

Out[146]: Text(0, 0.5, 'Final grade')



- Most students have a large family meaning there are more than 3 members in their family.
- Almost all the students live in the city with only very few in rural areas.
- Most students have a travel time to school less than 15 mins. Travel time and the number of students under that category are inversely proportional.
- Most students study between 0 to 5 hours. Students studying more than this are in very few numbers.
- Majority of the students do not have past class failures. Approximately 10% of students have past class failures.
- A lot of students have not paid for any extra classes but there is also significant number of students who have paid for it.
- Students do take part extra curricular activities but an almost 40% of students do not take part in any such activities.
- Almost all students have internet access at their homes with only a few students without such access.
- Most students were not in a romantic relationship, approximately 25% students were in one.
- Students have a really good relation with their family with only less than 10% students have a poor relation with the family
- The students also had a good amount of freetime.
- Most students did go out with their friends with only a small percentage.
- There is negative relation between final grade and number of absences. As the days of absence increases, the final grade decreases.

## Question 2

### a) Load the dataset

```
In #Load the dataset  
[147]: female = pd.read_csv("female_stud.csv")  
female
```

Out[147]:

	large_family	lives_in_city	traveltime	studytime	failures	pai
0	1	1	2	2	0	0
1	1	1	1	2	0	0
2	0	1	1	2	3	1
3	1	1	1	3	0	1
4	1	1	1	2	0	1
...	...	...	...	...	...	...
203	1	0	2	3	0	1
204	1	0	3	1	0	1
205	1	0	1	3	1	0
206	0	1	1	2	0	1
207	1	1	2	2	1	0

208 rows × 14 columns



### b) Inspect the data

```
In #Dimension of the dataset  
[148]: female.shape
```

Out[148]: (208, 14)

There are 208 students and 14 indicators included in the dataset.

```
In [149]: #Comparing indicators of male and female dataset  
print(list(male.columns)==list(female.columns))  
list(female.columns)
```

True

```
Out[149]: ['large_family',  
           'lives_in_city',  
           'travelttime',  
           'studytime',  
           'failures',  
           'paid',  
           'activities',  
           'internet',  
           'romantic',  
           'famrel',  
           'freetime',  
           'goout',  
           'absences',  
           'final_grade']
```

### c) Perform t-test

To check whether any indicators differ in the male and female group, we perform the t-test. We can consider 2 hypothesis:

- H0 - The indicators are equal in the male and female group
- H1 - The indicators differ in the male and female group

The significance level is taken as 0.01

```
In [150]: #Perform t-test on the indicators
alpha = 0.01
t_values, p_values = stats.ttest_ind(male, female, equal_var=True)
ttest_df = pd.DataFrame({"Indicators": male.columns,
                           "t-statistic": t_values,
                           "p-value": p_values})

#Comparing p-value with significance level
print(ttest_df)
ttest_df.loc[ttest_df['p-value'] < alpha, 'Indicators']
```

	Indicators	t-statistic	p-value
0	large_family	-1.788677	7.443690e-02
1	lives_in_city	-0.565304	5.721898e-01
2	traveltime	1.186055	2.363171e-01
3	studytime	-6.378011	5.045044e-10
4	failures	0.881778	3.784358e-01
5	paid	-2.581427	1.020069e-02
6	activities	1.989053	4.738832e-02
7	internet	0.875356	3.819147e-01
8	romantic	-2.033136	4.271032e-02
9	famrel	1.171091	2.422716e-01
10	freetime	4.873860	1.590705e-06
11	goout	1.508960	1.321126e-01
12	absences	-1.330451	1.841411e-01
13	final_grade	2.061993	3.986533e-02

```
Out[150]: 3    studytime
          10   freetime
Name: Indicators, dtype: object
```

From the t-test we see that, p-values of **studytime** and **freetime** are lesser than 0.01. Therefore we reject the null hypothesis for these two and accept it for the rest of the indicators. This means that there is a significant difference in the indicators **studytime** and **freetime** for male and female groups. When the p-value is lesser than the significance level it means that the result has not happened due to a random chance.

The indicators used in the female dataset are the same as that of the male dataset.

## Question 3

### a) Combine the dataframes

```
In #Combine the dataset vertically
[151]: student = pd.concat([male, female], axis=0)
student.reset_index(inplace=True, drop=True)
student
```

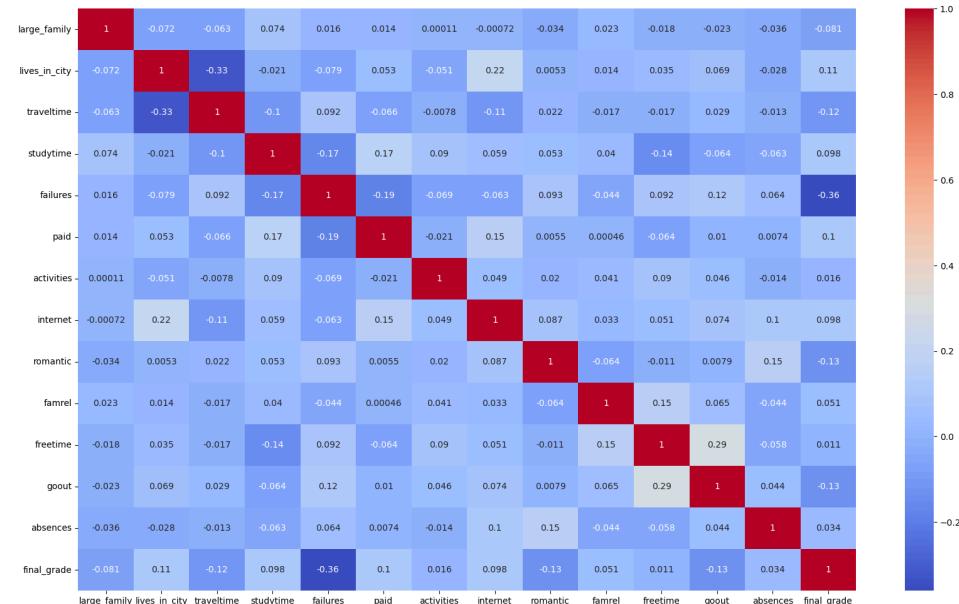
Out[151]:		large_family	lives_in_city	traveltime	studytime	failures	pai
	0	0	1	1	2	0	1
	1	0	1	1	2	0	0
	2	0	1	1	2	0	1
	3	1	1	1	2	0	1
	4	0	1	1	1	0	1
	...	...	...	...	...	...	...
	390	1	0	2	3	0	1
	391	1	0	3	1	0	1
	392	1	0	1	3	1	0
	393	0	1	1	2	0	1
	394	1	1	2	2	1	0

395 rows × 14 columns



## b) Compute Pearson correlation coefficient

```
In [152]: #Compute Pearson correlation coefficient for every indicator
plt.figure(figsize=(20,12))
corr_mat = student.corr(method='pearson')
sns.heatmap(corr_mat, annot=True, cmap='coolwarm')
plt.show()
```



Find the most correlated pairs

```
In [153]: #Find top 4 correlated pairs
corr_mat_abs = corr_mat.abs()

#Convert to upper triangular form
upper_mat = corr_mat_abs.where(np.triu(np.ones(corr_mat_abs.shape),
k=1).astype(bool))

#Drop all the null values
unique_pairs = upper_mat.unstack().dropna()

#Sort the values
sorted_mat = unique_pairs.sort_values()
print(sorted_mat)
```

```
activities    large_family    0.000113
famrel        paid            0.000460
internet      large_family    0.000720
romantic      lives_in_city  0.005257
                    paid          0.005536
                           ...
paid           failures        0.188039
internet      lives_in_city  0.216842
goout          freetime        0.285019
traveltime    lives_in_city  0.328096
final_grade   failures        0.360415
Length: 91, dtype: float64
```

The four most strongly correlated pairs are:

- final\_grade and failures
- traveltime and lives\_in\_city
- goout and freetime
- internet and lives\_in\_city

### c) Scatterplots

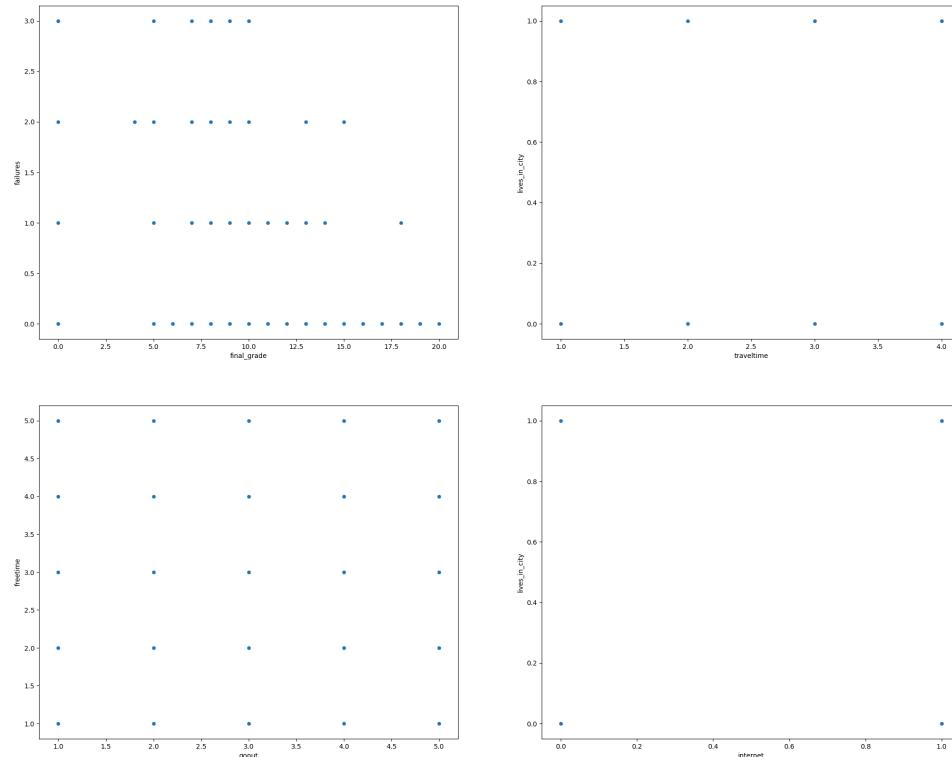
```
In [154]: #Scatter plots for correlated pairs
fig = plt.figure(figsize=(25,20))
plt.subplot(2,2,1)
sns.scatterplot(x='final_grade', y='failures', data=student)

plt.subplot(2,2,2)
sns.scatterplot(x='traveltime', y='lives_in_city', data=student)

plt.subplot(2,2,3)
sns.scatterplot(x='goout', y='freetime', data=student)

plt.subplot(2,2,4)
sns.scatterplot(x='internet', y='lives_in_city', data=student)
```

Out[154]: <Axes: xlabel='internet', ylabel='lives\_in\_city'>



From the plot between final grade and failures, it can be easily seen that there is negative correlation. For the other plots, the correlation is harder to understand as most of the variables are like categorical values.

## Question 4

### a) Create new column

```
In #Create new column Result  
[155]: student['Result'] = (student['final_grade']>9.5).astype(int)  
student.head()
```

	large_family	lives_in_city	traveltime	studytime	failures	paid
0	0	1	1	2	0	1
1	0	1	1	2	0	0
2	0	1	1	2	0	1
3	1	1	1	2	0	1
4	0	1	1	1	0	1



A new column **Result** is created which indicates the pass or fail status of student. The column has binary values

- 0 - The student has failed
- 1 - The student has passed

### b) Split data and standardise the variables

```
In #Response variable  
[156]: student_response = student['Result']  
#Predictor variables  
student_predictor = student.drop(['Result','final_grade'], axis=1)  
#Standardise the variables  
student_predictor_std = (student_predictor-  
student_predictor.mean())/student_predictor.std()
```

### c) Fit logistic regression model

```
In [157]: #Import model  
from sklearn.linear_model import LogisticRegression  
type(student_response)  
import statsmodels.api as sm
```

```
In [158]: student_predictor_std.insert(0,'intercept',1)  
student_predictor_std
```

Out[158]:		intercept	large_family	lives_in_city	traveltime	studytime	
0	1	-1.568015	0.534714	-0.642435	-0.042232	-0	
1	1	-1.568015	0.534714	-0.642435	-0.042232	-0	
2	1	-1.568015	0.534714	-0.642435	-0.042232	-0	
3	1	0.636134	0.534714	-0.642435	-0.042232	-0	
4	1	-1.568015	0.534714	-0.642435	-1.233786	-0	
...	...	...	...	...	...	...	...
390	1	0.636134	-1.865423	0.791247	1.149321	-0	
391	1	0.636134	-1.865423	2.224929	-1.233786	-0	
392	1	0.636134	-1.865423	-0.642435	1.149321	0	
393	1	-1.568015	0.534714	-0.642435	-0.042232	-0	
394	1	0.636134	0.534714	0.791247	-0.042232	0	

395 rows × 14 columns



```
In [159]: #Fit the model
logit1 = sm.Logit(student_response, student_predictor_std).fit()
print(logit1.summary())
```

```
Optimization terminated successfully.
    Current function value: 0.556201
    Iterations 5
                                         Logit Regression Results
=====
=====
Dep. Variable:                      Result      No. Observations: 395
Model:                             Logit      Df Residuals: 381
Method:                            MLE       Df Model: 13
Date:           Sat, 09 Dec 2023   Pseudo R-squ.: 0.1221
Time:           11:55:04          Log-Likelihood: -219.70
converged:                    True      LL-Null: -250.25
Covariance Type:                nonrobust   LLR p-value: 3.330e-08
=====
=====
                                         coef      std err      z      P>|z|
[0.025      0.975]
-----
-----
intercept      0.7801      0.118      6.602      0.000
0.548        1.012
large_family   -0.1204      0.121     -0.998      0.318
-0.357        0.116
lives_in_city  0.0507      0.125      0.404      0.686
-0.195        0.297
traveltime     0.0063      0.126      0.050      0.960
-0.241        0.254
studytime      0.0400      0.125      0.320      0.749
-0.205        0.285
failures       -0.6692      0.130     -5.132      0.000
-0.925        -0.414
paid            0.0768      0.121      0.633      0.527
-0.161        0.315
activities     -0.0269      0.119     -0.226      0.821
-0.260        0.206
internet       0.1361      0.119      1.147      0.251
-0.096        0.369
romantic       -0.1623      0.117     -1.383      0.167
-0.392        0.068
famrel          0.0644      0.118      0.547      0.584
-0.166        0.295
freetime         0.1361      0.126      1.079      0.280
-0.111        0.383
goout           -0.4254      0.127     -3.344      0.001
-0.675        -0.176
```

absences	-0.1279	0.111	-1.149	0.250
-0.346	0.090			
=====				
=====				

In logit1.aic  
[160]:  
Out[160]: 467.39887281774656

#### d) Forward selection using AIC

We can implement a forward selection for regression in the following way:

```
In #Function for forward AIC
[161]: def forwardAIC_logistic(X_std, Y, colnames):
    variables = ['intercept']
    mod = sm.Logit(Y,X_std[variables])
    res = mod.fit()
    mod_aic = res.aic
    aic_vals = dict()
    while(len(colnames)!=0):
        for p in colnames:
            variables.append(p)
            mod_new = sm.Logit(Y,X_std[variables])
            res_new = mod_new.fit()
            mod_aic_new = res_new.aic
            aic_vals[p] = mod_aic_new
            variables.remove(p)
        min_aic = min(aic_vals, key=aic_vals.get)
        min_aic_num = min(aic_vals)
        if min_aic in variables:
            break
        variables.append(min_aic)
        colnames.remove(min_aic)
    print(variables)
    return res_new
```

Use this function for the given data

```
In [162]: columns = list(student_predictor_std.columns)
columns.remove('intercept')
logistic_new = forwardAIC_logistic(student_predictor_std,
student_response, columns)
```

```
Optimization terminated successfully.
    Current function value: 0.633549
    Iterations 4
Optimization terminated successfully.
    Current function value: 0.632665
    Iterations 5
Optimization terminated successfully.
    Current function value: 0.632203
    Iterations 4
Optimization terminated successfully.
    Current function value: 0.632577
    Iterations 4
Optimization terminated successfully.
    Current function value: 0.630725
    Iterations 5
Optimization terminated successfully.
    Current function value: 0.578633
    Iterations 5
Optimization terminated successfully.
    Current function value: 0.629230
    Iterations 5
Optimization terminated successfully.
    Current function value: 0.633472
    Iterations 4
Optimization terminated successfully.
    Current function value: 0.631684
    Iterations 4
Optimization terminated successfully.
    Current function value: 0.628837
    Iterations 5
Optimization terminated successfully.
    Current function value: 0.632470
    Iterations 4
Optimization terminated successfully.
    Current function value: 0.633381
    Iterations 4
Optimization terminated successfully.
    Current function value: 0.616589
    Iterations 5
Optimization terminated successfully.
    Current function value: 0.629482
    Iterations 5
Optimization terminated successfully.
    Current function value: 0.577872
    Iterations 5
Optimization terminated successfully.
    Current function value: 0.578297
    Iterations 5
Optimization terminated successfully.
    Current function value: 0.578530
    Iterations 5
```

```
Optimization terminated successfully.  
    Current function value: 0.578472  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.578130  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.578574  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.577727  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.576334  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.578124  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.578563  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.566845  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.576188  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.565885  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.566075  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.566759  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.566802  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.566159  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.566834  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.565287  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.564350  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.565976  
    Iterations 5  
Optimization terminated successfully.  
    Current function value: 0.564931  
    Iterations 5
```

```
Optimization terminated successfully.
    Current function value: 0.564698
    Iterations 5
['intercept', 'failures', 'goout']
```

```
In [163]: print(logistic_new.summary())
          print(logistic_new.aic)
```

Logit Regression Results					
		Result	No. Observations:		
Dep. Variable:	395		Logit	Df Residuals:	391
Model:	391		MLE	Df Model:	3
Method:	3			Pseudo R-squ.:	0.1087
Date:	Sat, 09 Dec 2023		11:55:04	Log-Likelihood:	-223.06
Time:				LL-Null:	-250.25
converged:		True		Covariance Type:	nonrobust
				p-value:	9.253e-12
coef std err z P> z  [0.025					
0.975]					
-----					
intercept	0.7592	0.116	6.531	0.000	0.531
failures	-0.7061	0.128	-5.535	0.000	-0.956
goout	-0.3501	0.117	-2.980	0.003	-0.580
absences	-0.1435	0.108	-1.330	0.183	-0.355
0.068					
-----					
454.11132408212734					

The new model obtained using forward stepwise regression has the variables **failures** , **goout** and **absences** .

The AIC value for model obtained from stepwise regeression is approximately 454.11 and for the previous model is 467.40. So the model obtained using the forward stepwise regeression is a better fit as the AIC value is lower.

## Question 5

### a) Split data into training and test sets

```
In #Splitting data into response and predictor variables  
[164]: X = student.drop(['Result','final_grade'], axis=1)  
Y = student['final_grade']
```

The data is split into approximately 70% training data and 30% test data.

```
In #Splitting the data into test and train  
[165]: train_size = 277  
np.random.seed(123)  
train_select = np.random.permutation(range(len(Y)))  
X_train = X.iloc[train_select[:train_size],:].reset_index(drop=True)  
X_test = X.iloc[train_select[train_size:]],:].reset_index(drop=True)  
y_train = Y.iloc[train_select[:train_size]].reset_index(drop=True)  
y_test = Y.iloc[train_select[train_size:]].reset_index(drop=True)
```

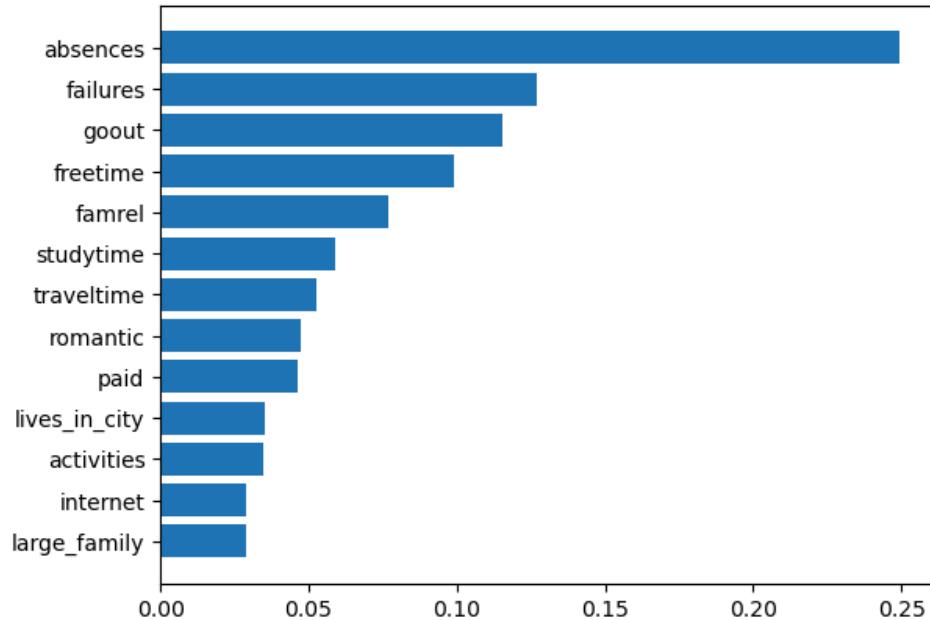
### b) Fit random forest regression model

```
In #Import random forest regressor  
[166]: from sklearn.ensemble import RandomForestRegressor  
#RandomForestRegressor?
```

```
In #Fitting the model  
[167]: rf = RandomForestRegressor(n_estimators=10, random_state=101)  
rf_fit = rf.fit(X_train, y_train)
```

```
In #Sorting important features
[168]: sorted_idx = rf.feature_importances_.argsort()
plt.barh(X.columns[sorted_idx], rf.feature_importances_[sorted_idx])
```

Out[168]: <BarContainer object of 13 artists>



The important variables for predicting the final grade of a student are:

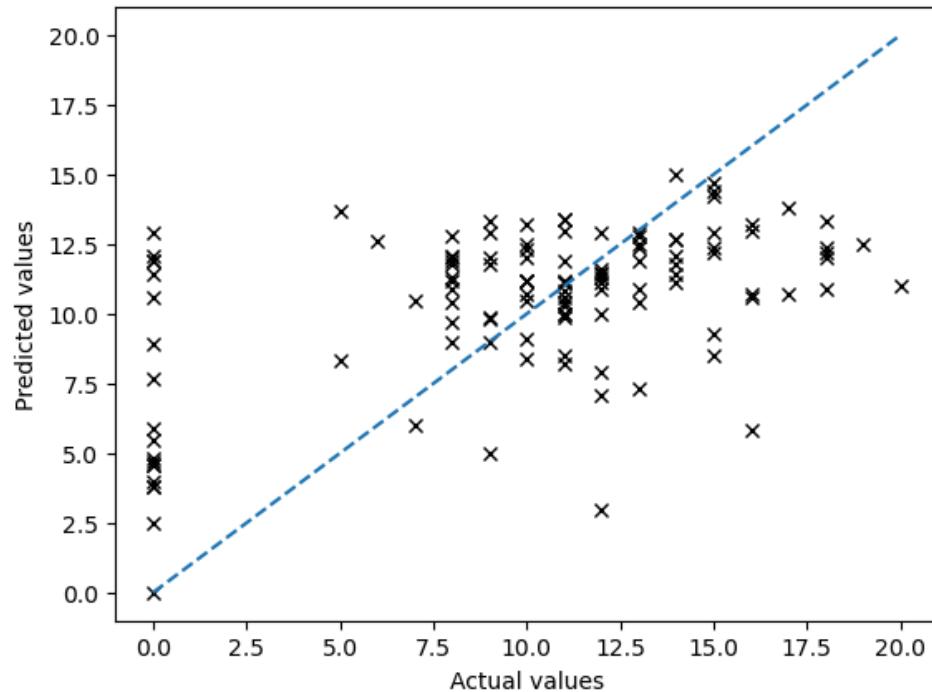
- absences
- failures
- goout
- freetime

The models fit in the question 4 and question 5 are very similar because the stepwise regression picked up the same features - absences, failures and goout. Stepwise regression does not give us an optimal model but it gives a better and good model. Now the random forest regressor model also suggests the same features to be important. So we can conclude that these features cannot be omitted for any other model as well.

### c) Predict and create scatter plot

```
In [169]: #Predict the values  
rf_test_pred = rf.predict(X_test)  
fig = plt.figure()  
plt.plot(y_test,rf_test_pred,'kx')  
plt.plot([0,20], [0,20], ls="--")  
plt.xlabel('Actual values')  
plt.ylabel('Predicted values')
```

Out[169]: Text(0, 0.5, 'Predicted values')



The model is a decent fit looking at the plot. There are outliers for the lower value predictions. For values 10 and above, the prediction is better as it is close to the regression line, but there are also some points spread away from the line indicating more error in the prediction.

#### d) Fit and predict using different models

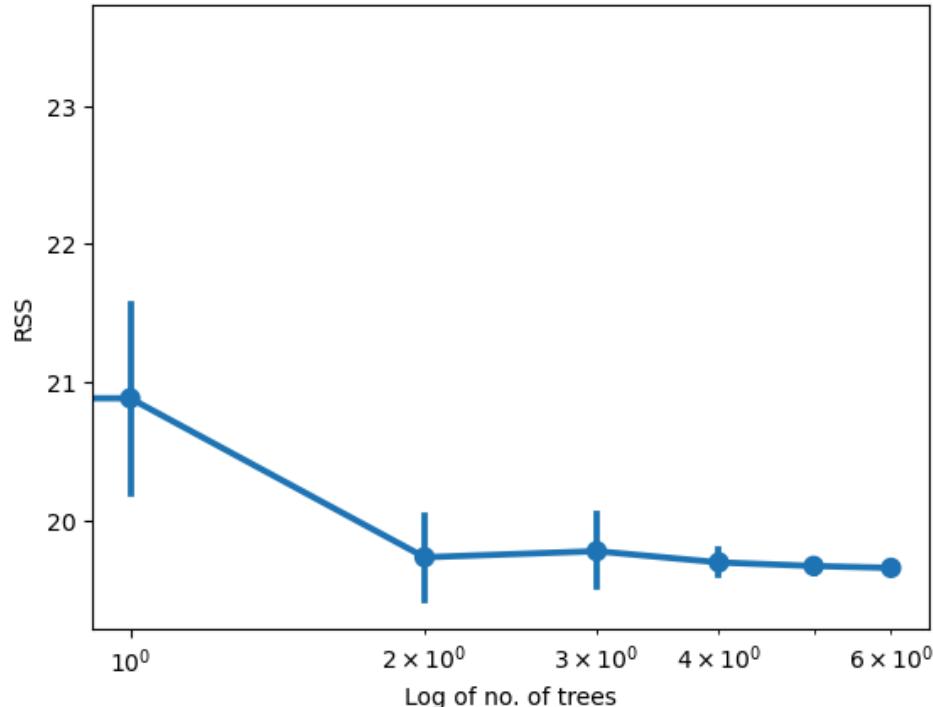
```
In [170]: #Fitting the model
trees = [5, 10, 50, 100, 500, 1000, 5000]
perf_df = pd.DataFrame(columns=['Trees','Performance'])
pred_values = []
for item in trees:
    for i in range(20):
        rf = RandomForestRegressor(n_estimators=item,
random_state=101+i)
        rf_fit = rf.fit(X_train, y_train)
        rf_test_pred = rf.predict(X_test)
        pred_values.append(rf_test_pred)
        RSS_rf = np.mean(pow((rf_test_pred- y_test),2))
        perf_list = [item, RSS_rf]
        perf_df.loc[len(perf_df)] = perf_list
perf_df
```

Out[170]:

	Trees	Performance
0	5.0	22.870847
1	5.0	23.177213
2	5.0	24.602524
3	5.0	26.756761
4	5.0	21.403729
...	...	...
135	5000.0	19.729556
136	5000.0	19.662193
137	5000.0	19.623193
138	5000.0	19.686962
139	5000.0	19.656232

140 rows × 2 columns

```
In [171]: #Plot the performance metric  
plt.scatter(perf_df['Trees'], perf_df['Performance'], alpha=0.5)  
sns.pointplot(x='Trees', y='Performance', data=perf_df)  
plt.xscale('log')  
plt.xlabel('Log of no. of trees')  
plt.ylabel('RSS')  
plt.show()
```



As the number of trees in the model increase, the RSS value and the standard deviation decreases. This means that the performance of the model is better but it also increases the computation time.

### e) Models with different random states

Random state controls the shuffling of the data. If we just perform one fit with the a particular random state, we may not get proper RSS value. When we fit the model with different random states multiple times, the average RSS value should be close to the true RSS value for the dataset.

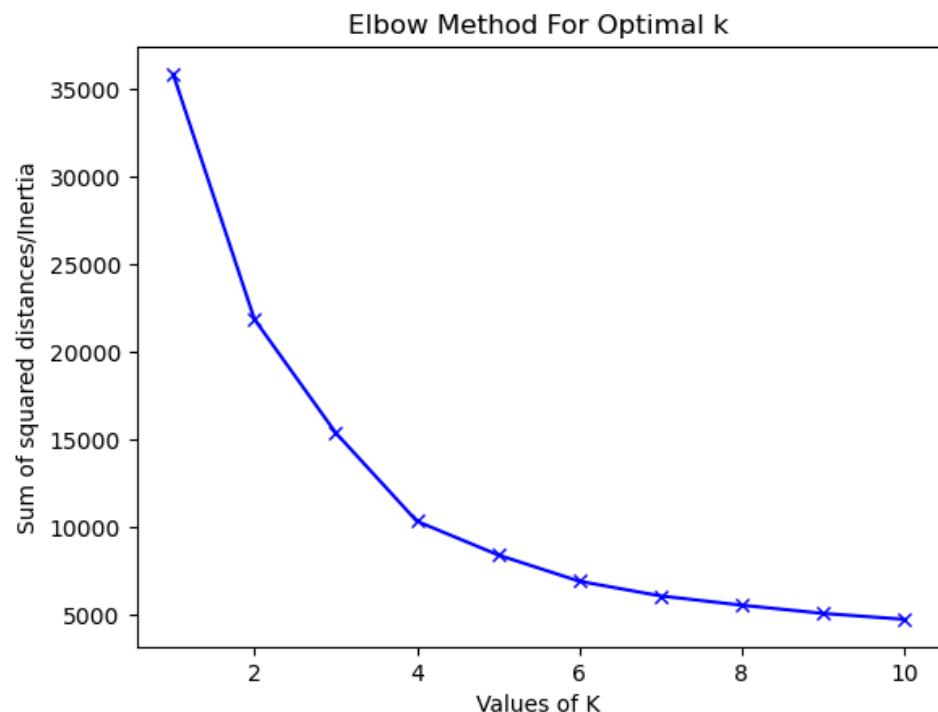
## Question 6

### a) K-means cluster analysis

```
In [172]: import warnings
warnings.filterwarnings('ignore')

In [173]: #Import model
from sklearn.cluster import KMeans

In [174]: #Fit k-means model
k_student = student.drop(['Result'], axis=1)
sse = []
for k in range(1, 11):
    km = KMeans(n_clusters=k, n_init=10)
    km.fit(k_student)
    sse.append(km.inertia_)
plt.plot(range(1,11), sse, 'bx-')
plt.xlabel('Values of K')
plt.ylabel('Sum of squared distances/Inertia')
plt.title('Elbow Method For Optimal k')
plt.show()
```



The optimal value for k is 4.

## b) K-means cluster analysis with optimal k

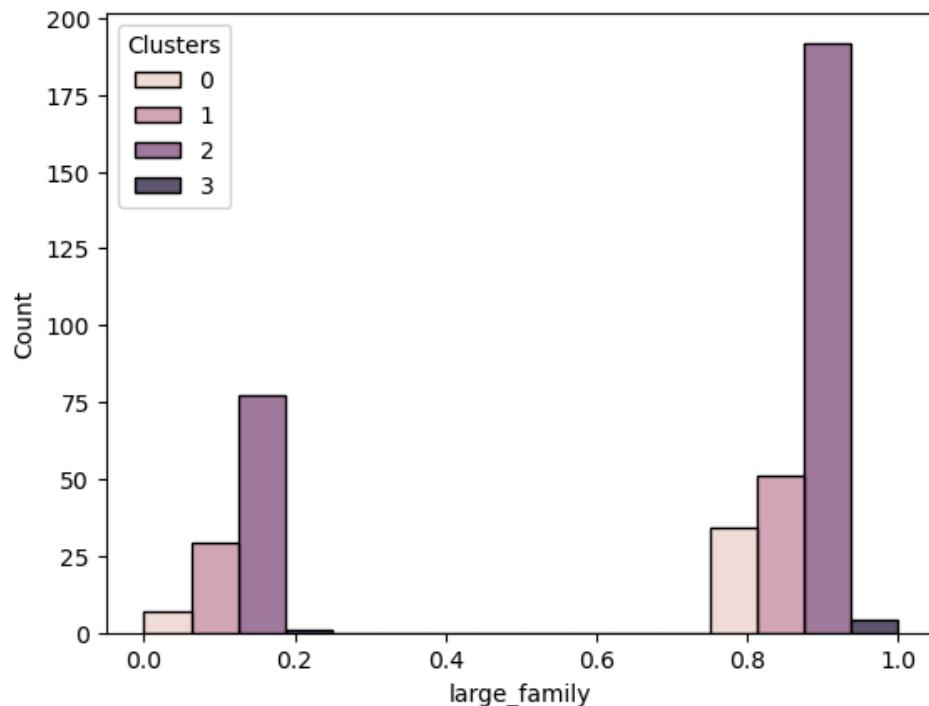
```
In [175]: km = KMeans(n_clusters=4, n_init=10)
km.fit(k_student)
preds = km.predict(k_student)
```

```
In [176]: type(preds)
Out[176]: numpy.ndarray
```

```
In [177]: #sns.countplot(x=preds, hue=k_student['large_family'], stat='percent')
pred_df = pd.DataFrame()
pred_df['Clusters'] = preds
temp_student = pd.concat([k_student, pred_df], axis='columns')
filter_df0 = temp_student[temp_student['Clusters']==0]
filter_df1 = temp_student[temp_student['Clusters']==1]
filter_df2 = temp_student[temp_student['Clusters']==2]
filter_df3 = temp_student[temp_student['Clusters']==3]
```

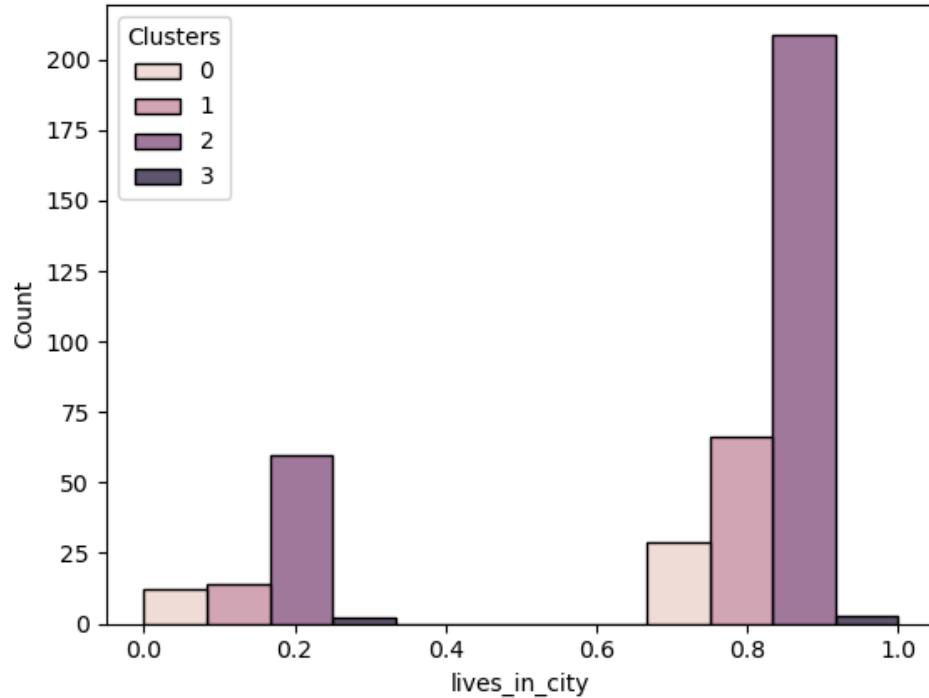
```
In [178]: sns.histplot(x = temp_student['large_family'],
hue=temp_student['Clusters'], multiple="dodge", bins=4)
```

Out[178]: <Axes: xlabel='large\_family', ylabel='Count'>



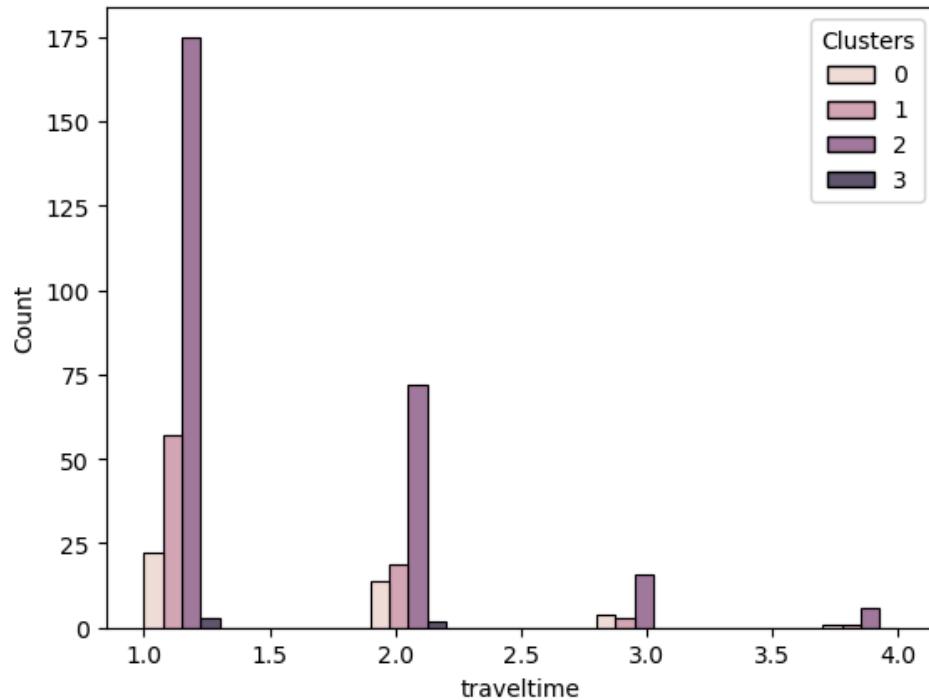
```
In [179]: sns.histplot(x = temp_student['lives_in_city'],
hue=temp_student['Clusters'], multiple="dodge", bins=3)
```

Out[179]: <Axes: xlabel='lives\_in\_city', ylabel='Count'>



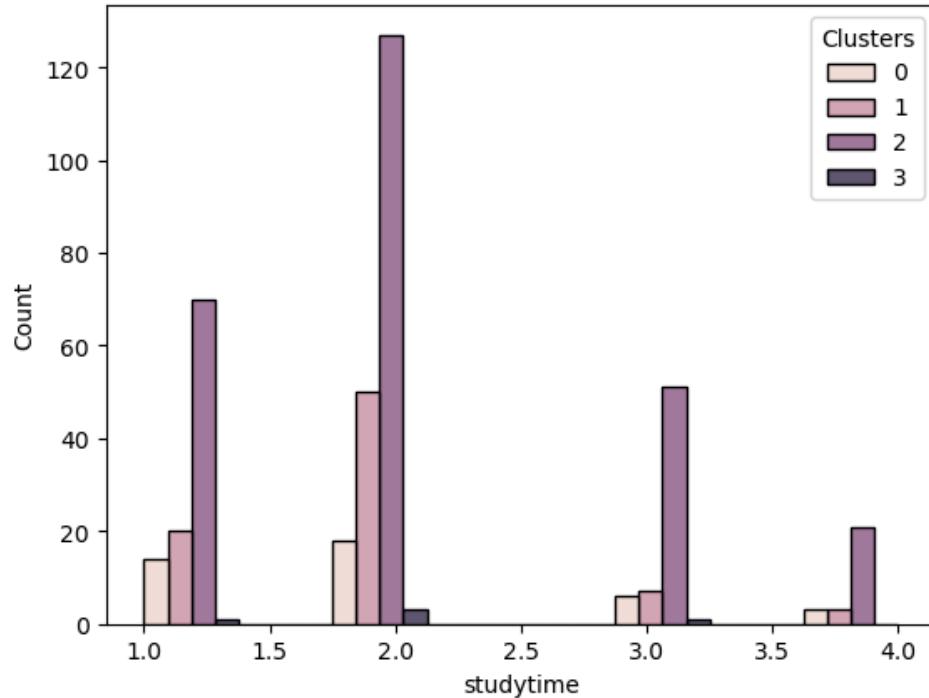
```
In [180]: sns.histplot(x = temp_student['traveltime'],
hue=temp_student['Clusters'], multiple="dodge", bins=10)
```

Out[180]: <Axes: xlabel='traveltime', ylabel='Count'>



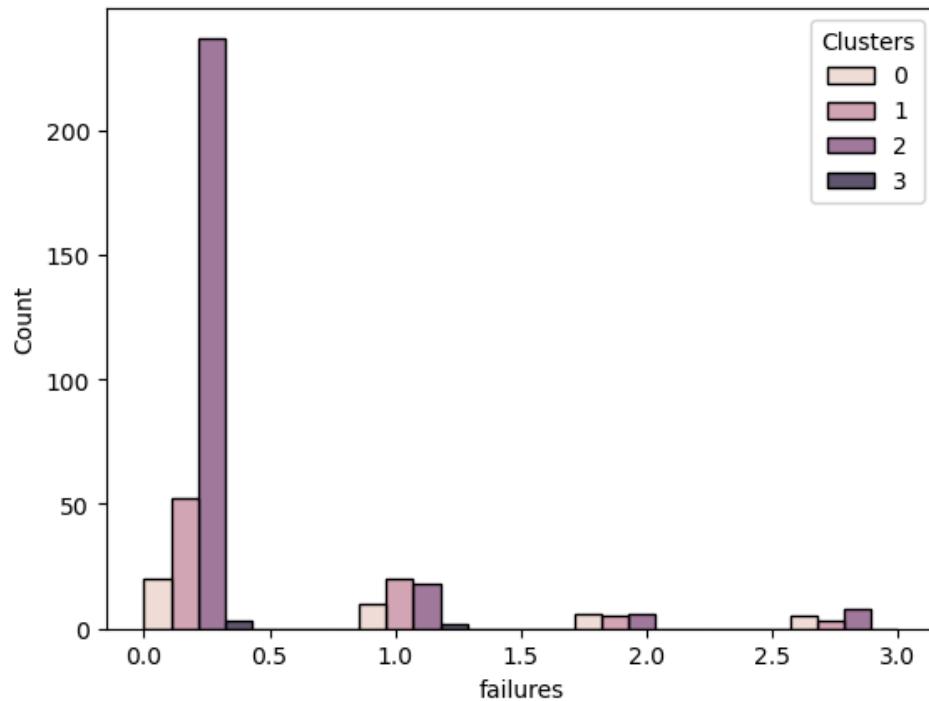
```
In [181]: sns.histplot(x = temp_student['studytime'],  
hue=temp_student['Clusters'], multiple="dodge", bins=8)
```

Out[181]: <Axes: xlabel='studytime', ylabel='Count'>



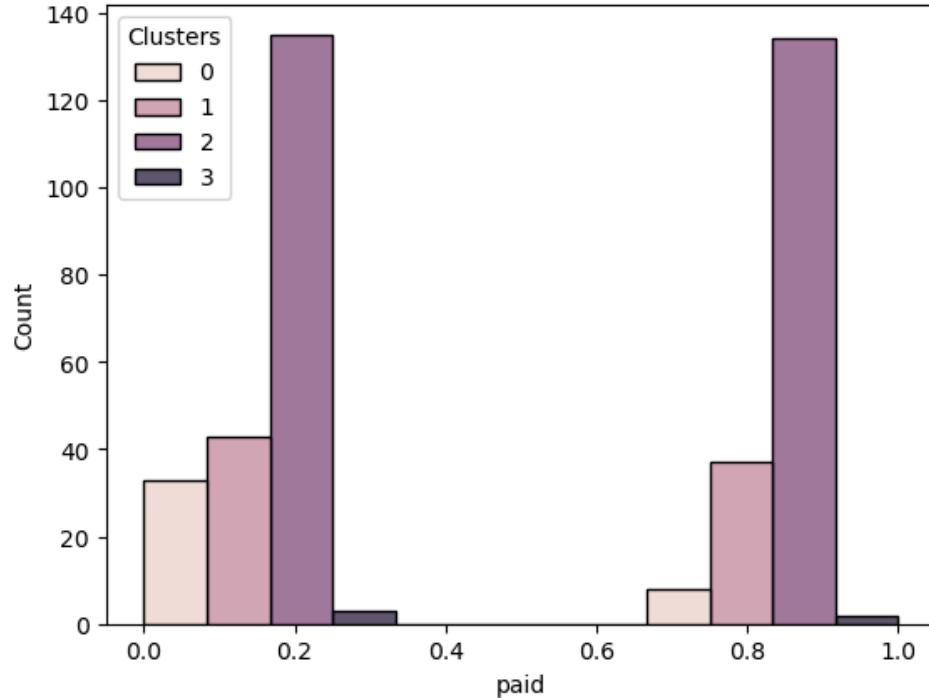
```
In [182]: sns.histplot(x = temp_student['failures'], hue=temp_student['Clusters'],  
multiple="dodge", bins=7)
```

Out[182]: <Axes: xlabel='failures', ylabel='Count'>



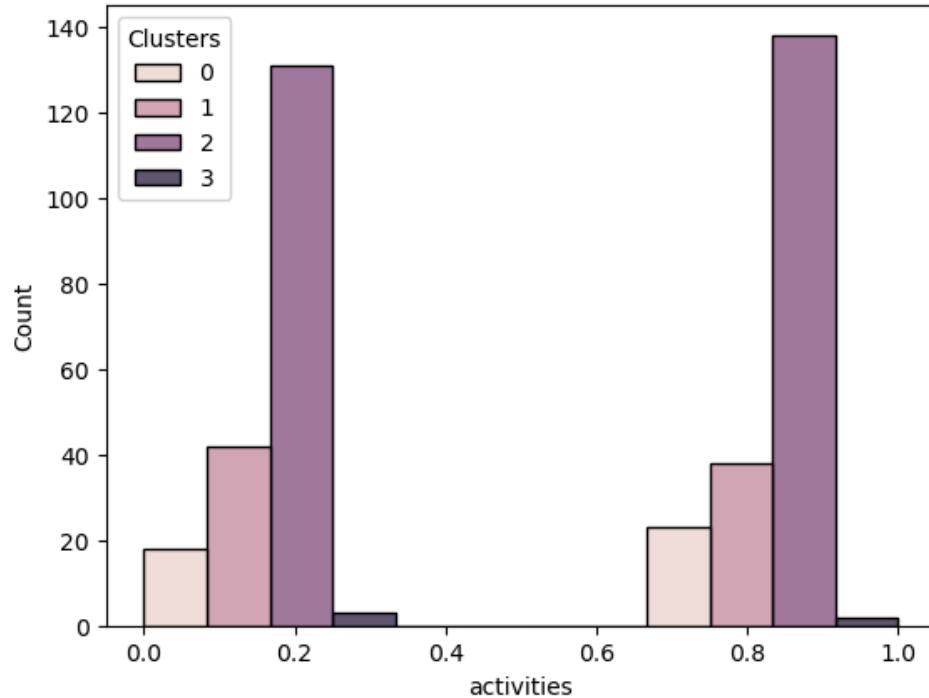
```
In [183]: sns.histplot(x = temp_student['paid'], hue=temp_student['Clusters'],  
multiple="dodge", bins=3)
```

Out[183]: <Axes: xlabel='paid', ylabel='Count'>



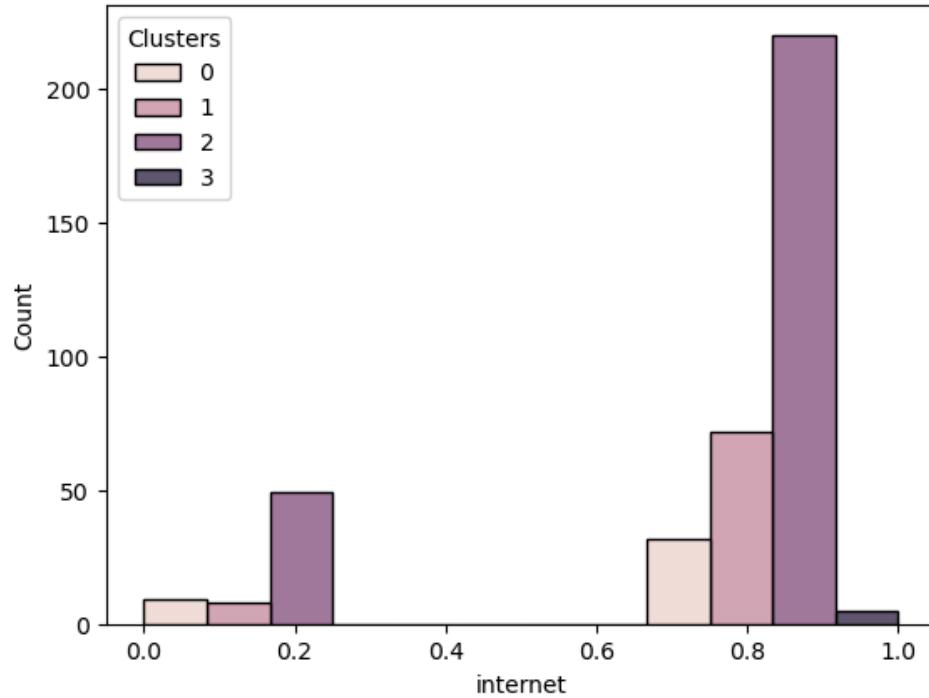
```
In [184]: sns.histplot(x = temp_student['activities'],  
hue=temp_student['Clusters'], multiple="dodge", bins=3)
```

Out[184]: <Axes: xlabel='activities', ylabel='Count'>



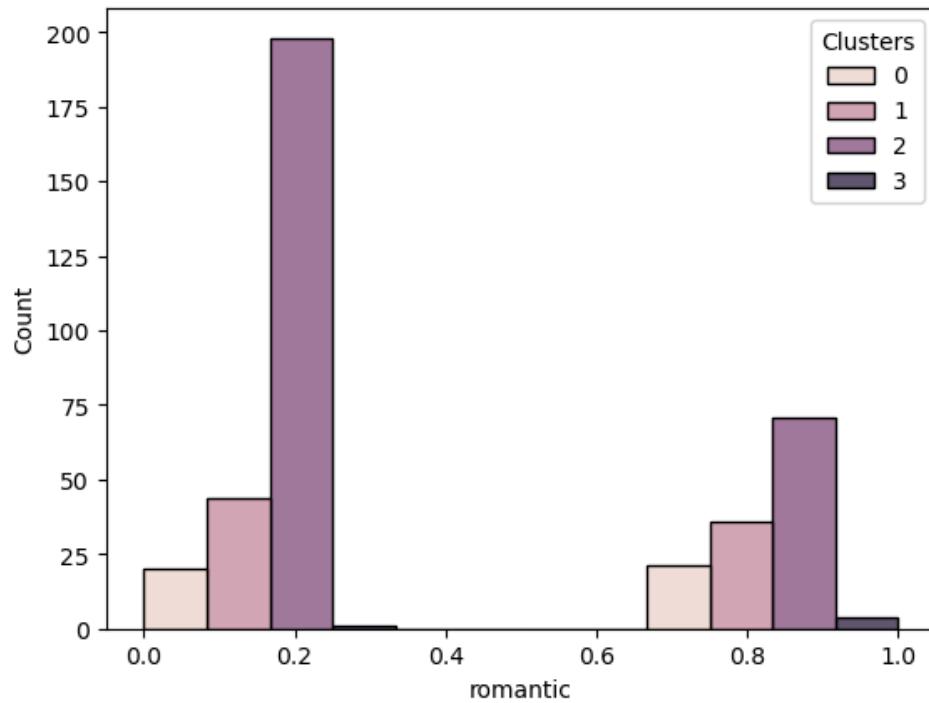
```
In [185]: sns.histplot(x = temp_student['internet'], hue=temp_student['Clusters'],  
multiple="dodge", bins=3)
```

Out[185]: <Axes: xlabel='internet', ylabel='Count'>



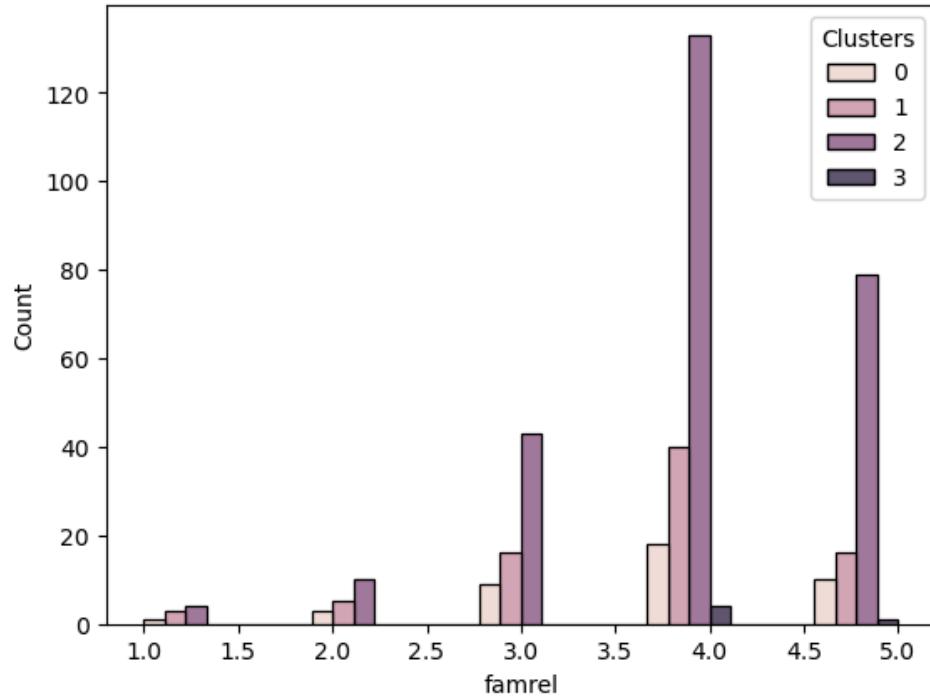
```
In [186]: sns.histplot(x = temp_student['romantic'], hue=temp_student['Clusters'],  
multiple="dodge", bins=3)
```

Out[186]: <Axes: xlabel='romantic', ylabel='Count'>



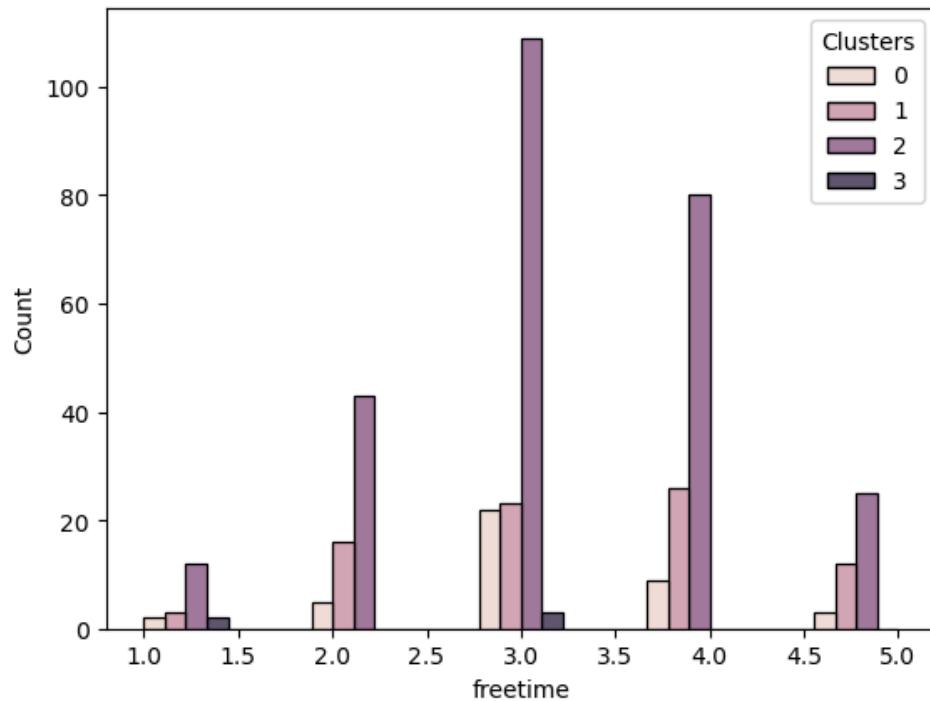
```
In [187]: sns.histplot(x = temp_student['famrel'], hue=temp_student['Clusters'],  
multiple="dodge", bins=9)
```

Out[187]: <Axes: xlabel='famrel', ylabel='Count'>



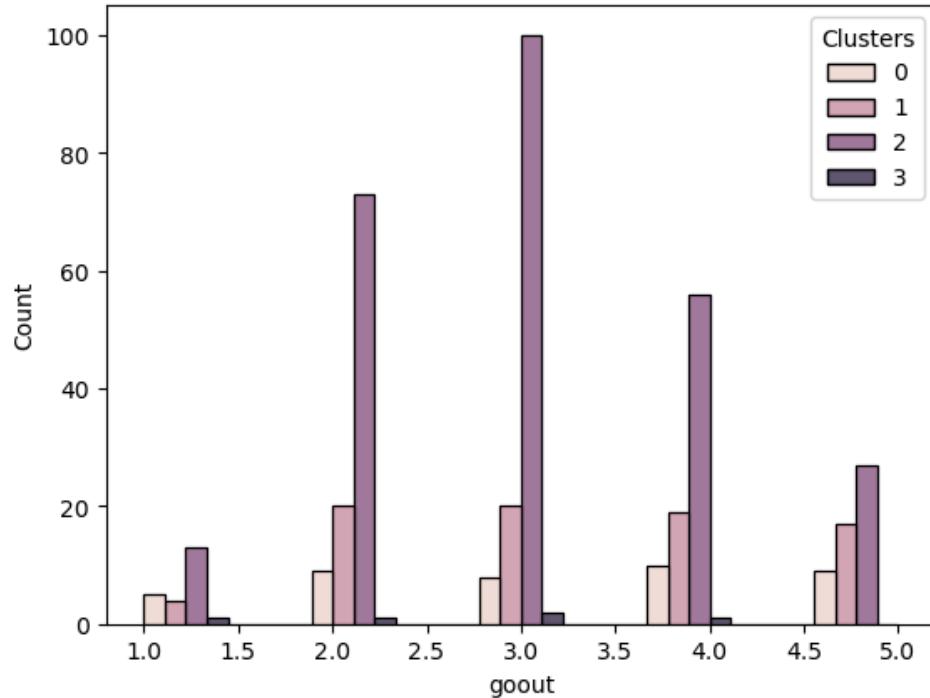
```
In [188]: sns.histplot(x = temp_student['freetime'], hue=temp_student['Clusters'],  
multiple="dodge", bins=9)
```

Out[188]: <Axes: xlabel='freetime', ylabel='Count'>



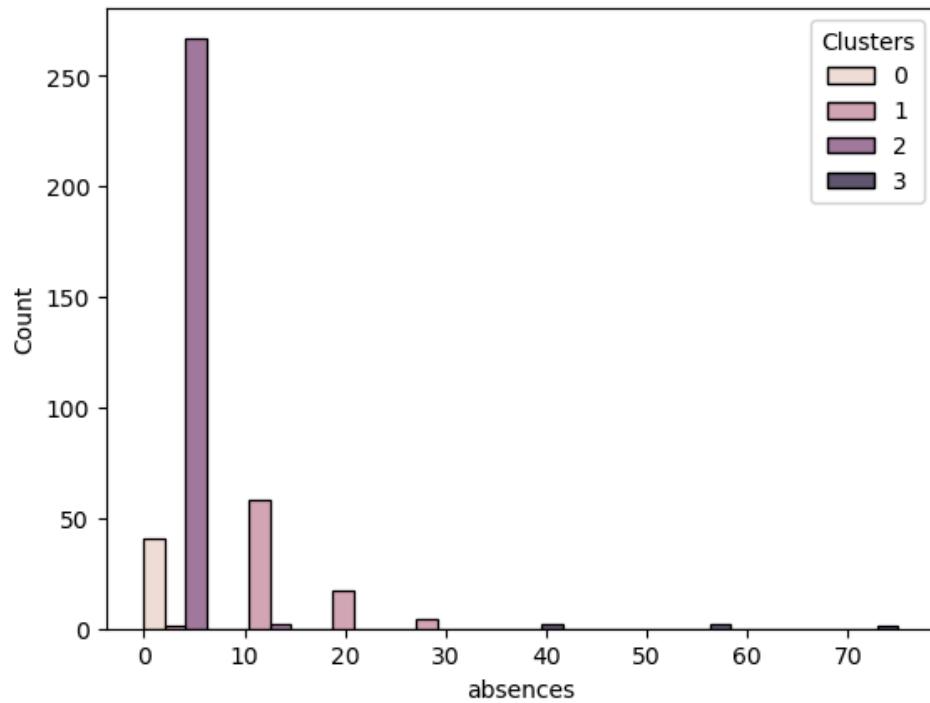
```
In [189]: sns.histplot(x = temp_student['goout'], hue=temp_student['Clusters'],  
multiple="dodge", bins=9)
```

Out[189]: <Axes: xlabel='goout', ylabel='Count'>



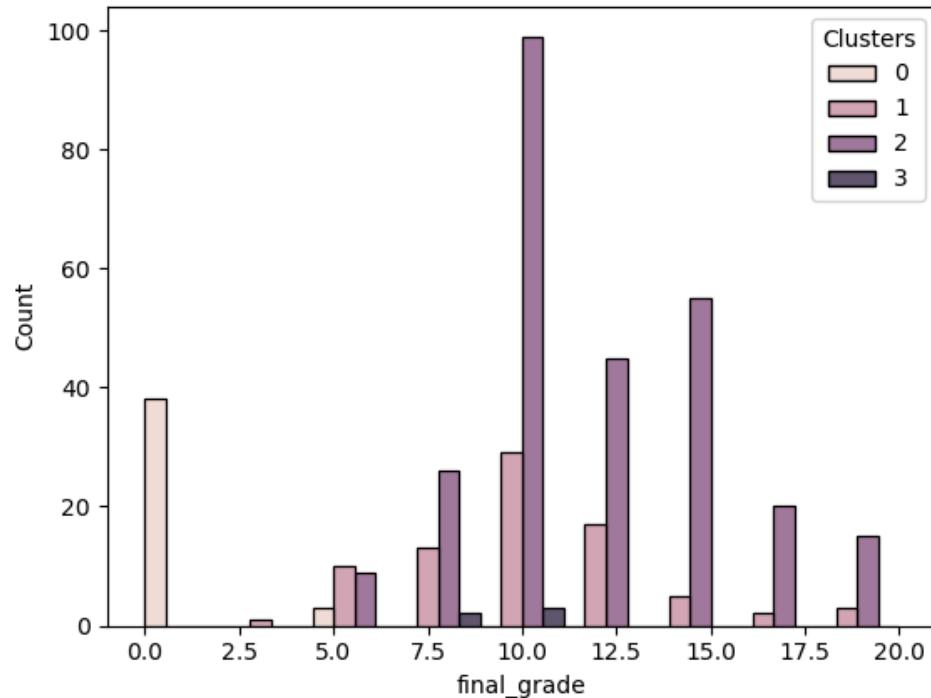
```
In [190]: sns.histplot(x = temp_student['absences'], hue=temp_student['Clusters'],  
multiple="dodge", bins=9)
```

Out[190]: <Axes: xlabel='absences', ylabel='Count'>



```
In [191]: sns.histplot(x = temp_student['final_grade'],  
hue=temp_student['Clusters'], multiple="dodge", bins=9)
```

Out[191]: <Axes: xlabel='final\_grade', ylabel='Count'>

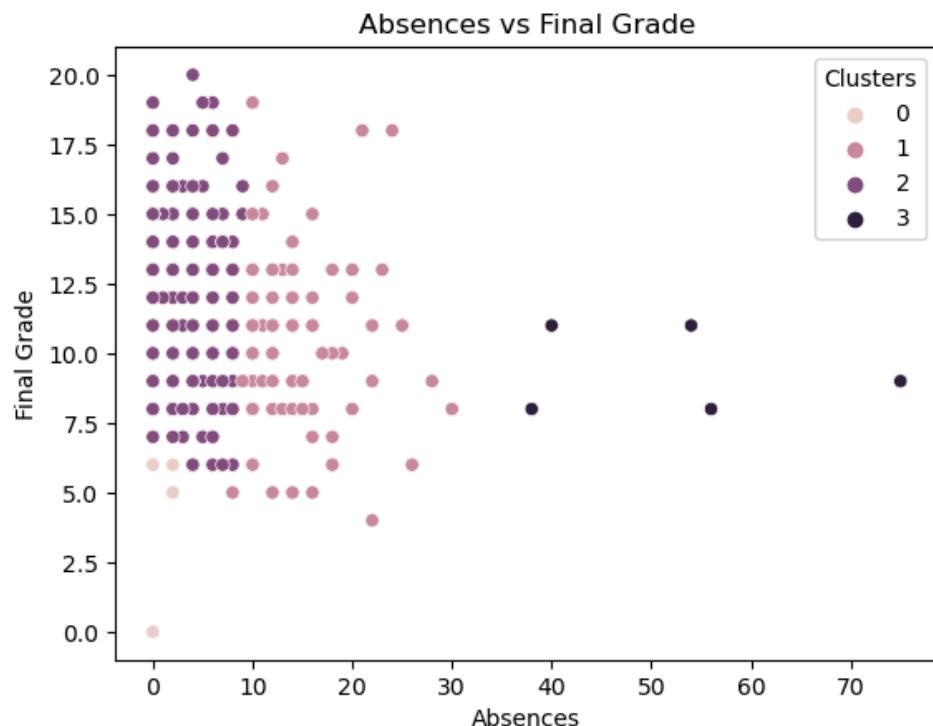


Discriminatory - final\_grade, absences, travelttime, failures(maybe)

### c) Scatter plots

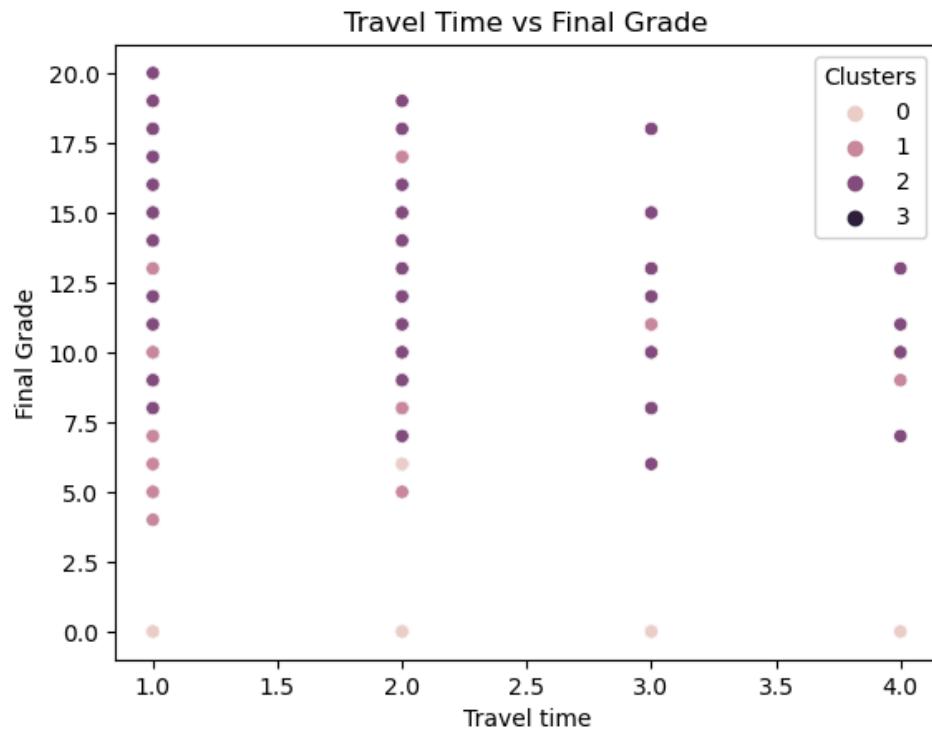
```
In [192]: #Scatter plot fpr absences and final grade  
sns.scatterplot(x = temp_student['absences'], y =  
temp_student['final_grade'], hue = temp_student['Clusters'])  
plt.xlabel('Absences')  
plt.ylabel('Final Grade')  
plt.title('Absences vs Final Grade')
```

Out[192]: Text(0.5, 1.0, 'Absences vs Final Grade')



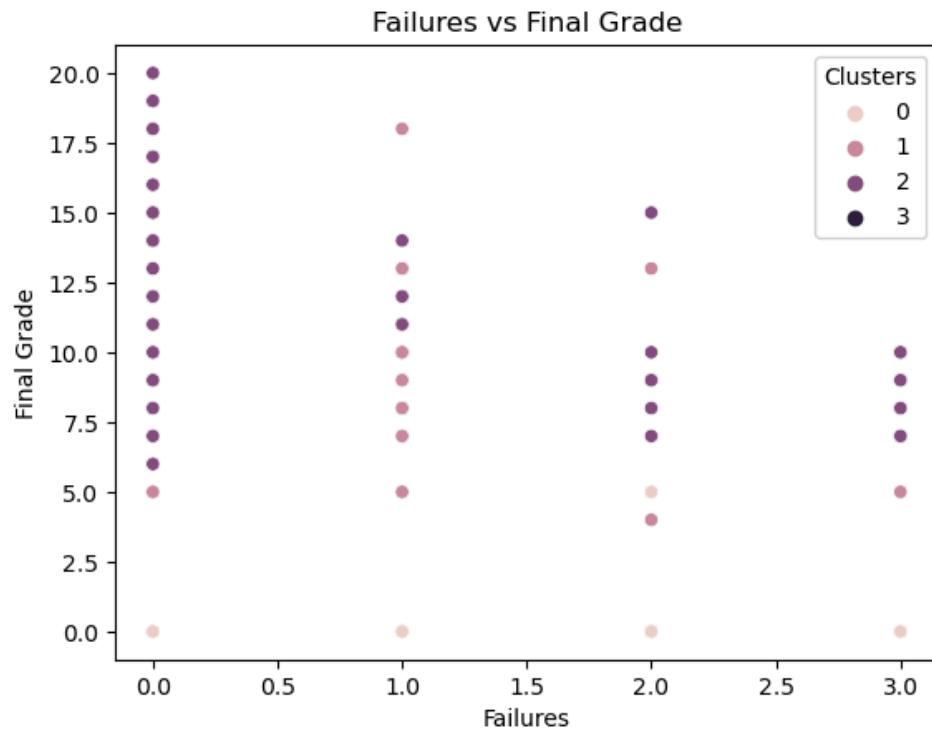
```
In [193]: #Scatter plot for travel time and final grade  
sns.scatterplot(x = temp_student['traveltime'], y =  
temp_student['final_grade'], hue = temp_student['Clusters'])  
plt.xlabel('Travel time')  
plt.ylabel('Final Grade')  
plt.title('Travel Time vs Final Grade')
```

Out[193]: Text(0.5, 1.0, 'Travel Time vs Final Grade')



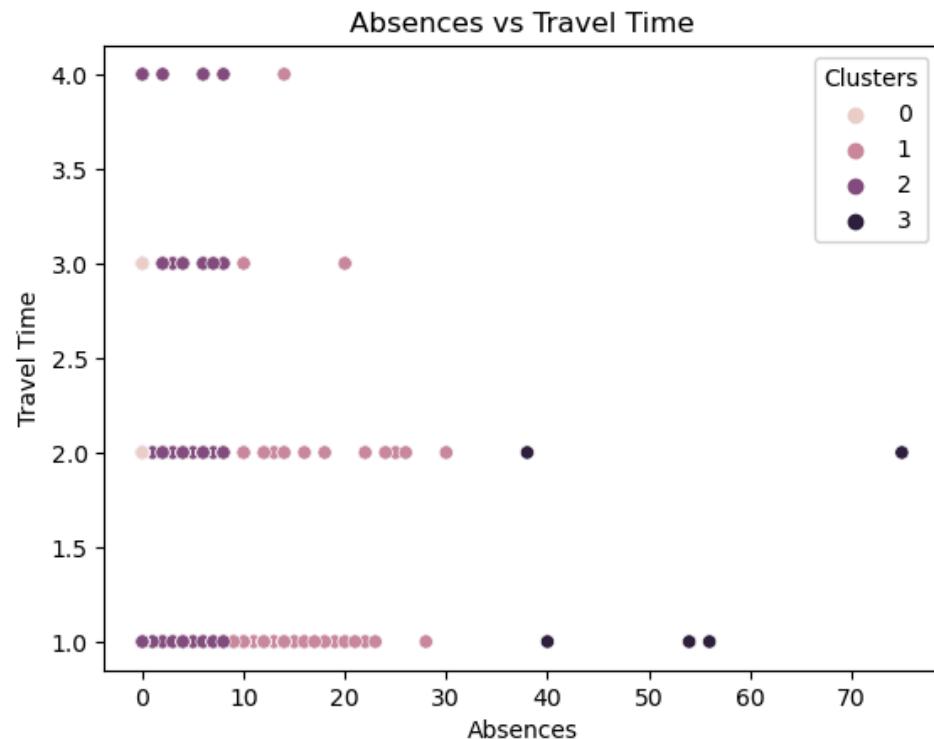
```
In [194]: #Scatter plot for failures and final grade  
sns.scatterplot(x = temp_student['failures'], y =  
temp_student['final_grade'], hue = temp_student['Clusters'])  
plt.xlabel('Failures')  
plt.ylabel('Final Grade')  
plt.title('Failures vs Final Grade')
```

Out[194]: Text(0.5, 1.0, 'Failures vs Final Grade')



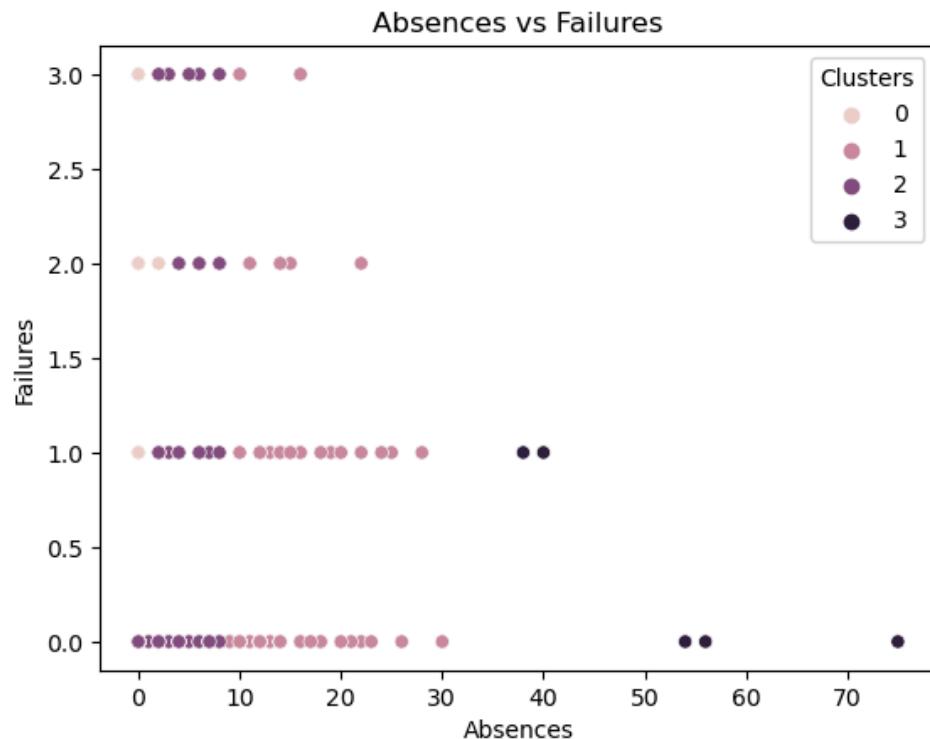
```
In [195]: #Scatter plot for absences and travel time  
sns.scatterplot(x = temp_student['absences'], y =  
temp_student['traveltime'], hue = temp_student['Clusters'])  
plt.xlabel('Absences')  
plt.ylabel('Travel Time')  
plt.title('Absences vs Travel Time')
```

Out[195]: Text(0.5, 1.0, 'Absences vs Travel Time')



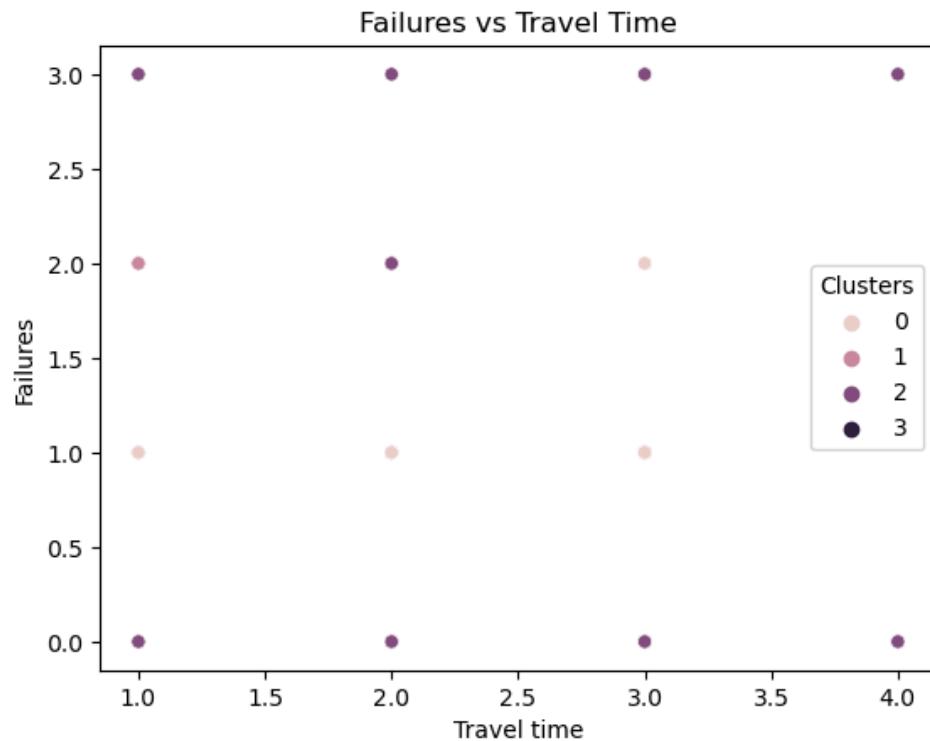
```
In #Scatter plot for absences and failures  
[196]: sns.scatterplot(x = temp_student['absences'], y =  
temp_student['failures'], hue = temp_student['Clusters'])  
plt.xlabel('Absences')  
plt.ylabel('Failures')  
plt.title('Absences vs Failures')
```

Out[196]: Text(0.5, 1.0, 'Absences vs Failures')



```
In #Scatter plot for travel time and failures  
[197]: sns.scatterplot(x = temp_student['traveltime'], y =  
temp_student['failures'], hue = temp_student['Clusters'])  
plt.xlabel('Travel time')  
plt.ylabel('Failures')  
plt.title('Failures vs Travel Time')
```

Out[197]: Text(0.5, 1.0, 'Failures vs Travel Time')



From the scatter plots it can be seen that the clusters are being created for certain ranges of values of the variables. In some cases it is more evident, like in "Absences vs Failures", "Absences vs Travel Time" and "Absences vs Final Grade". From the plots "Absences" seem to be the most discriminatory variable for the clustering. This means that there are students with some distinct characteristics which can be mapped to different classes.

## d) Different clustering algorithm - DBSCAN

DBSCAN (Density-Based Spatial Clustering of Applications with Noise) is a clustering algorithm that categorises the data into clusters based on their density. It does not require a cluster count to be given. It can also find outliers and classify it into noise instead of it being in a cluster. It can find clusters of arbitrary shapes unlike in K-means.

It has main two parameters, **Epsilon** and **Minimum points**. Epsilon is the maximum distance to be considered between two neighbors and Minimum points is the minimum number of points required to form a dense region.

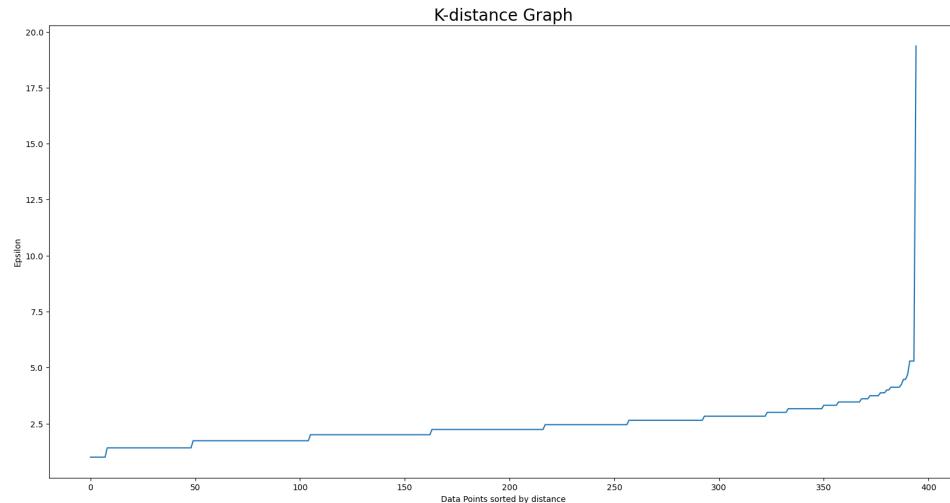
This algorithm is best used when the clusters are irregularly shaped and there are outliers in the dataset

```
In [200]: #Import the model
          from sklearn.cluster import DBSCAN
```

We plot a k-distance to find the optimal Epsilon value. The optimal value of Epsilon is the point where there is maximum curvature in the graph.

```
In [201]: #Fit nearest neighbors
          from sklearn.neighbors import NearestNeighbors
          neigh = NearestNeighbors(n_neighbors=2)
          nbrs = neigh.fit(k_student)
          distances, indices = nbrs.kneighbors(k_student)
```

```
In # Plot K-distance Graph
[202]: distances = np.sort(distances, axis=0)
         distances = distances[:,1]
         plt.figure(figsize=(20,10))
         plt.plot(distances)
         plt.title('K-distance Graph', fontsize=20)
         plt.xlabel('Data Points sorted by distance')
         plt.ylabel('Epsilon')
         plt.show()
```



The optimal Epsilon value is 3

```
In #Fit DBSCAN model
[203]: dbSCAN=DBSCAN(eps=3, min_samples=6)
         dbSCAN.fit(k_student)
         dbSCAN.labels_
         labels = dbSCAN.labels_
         n_clusters_ = len(set(labels)) - (1 if -1 in labels else 0)
         n_noise_ = list(labels).count(-1)

         print("Estimated number of clusters: %d" % n_clusters_)
         print("Estimated number of noise points: %d" % n_noise_)
```

```
Estimated number of clusters: 3
Estimated number of noise points: 87
```

## Plot histograms

```
In [204]: pred_df1 = pd.DataFrame()
pred_df1['Clusters'] = labels
temp_student1 = pd.concat([k_student, pred_df1], axis='columns')
#Removing noise
temp_student1 = temp_student1[temp_student1['Clusters'] != -1]
temp_student1
```

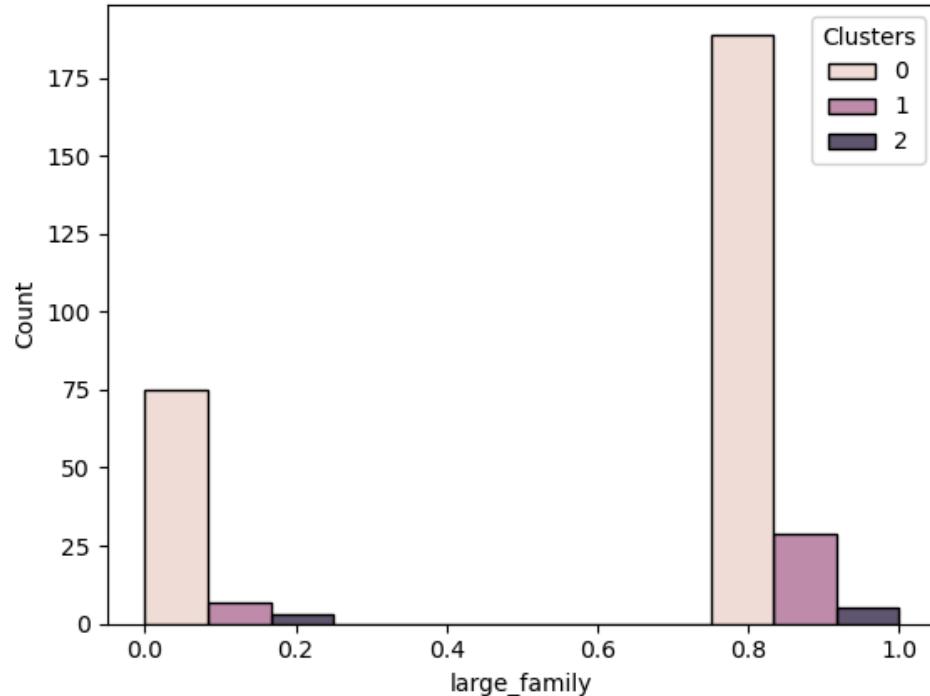
Out[204]:		large_family	lives_in_city	traveltime	studytime	failures	pai
	0	0	1	1	2	0	1
	1	0	1	1	2	0	0
	3	1	1	1	2	0	1
	4	0	1	1	1	0	1
	5	1	1	2	2	0	1
	...	...	...	...	...	...	...
	388	1	1	1	2	0	1
	390	1	0	2	3	0	1
	391	1	0	3	1	0	1
	392	1	0	1	3	1	0
	393	0	1	1	2	0	1

308 rows × 15 columns



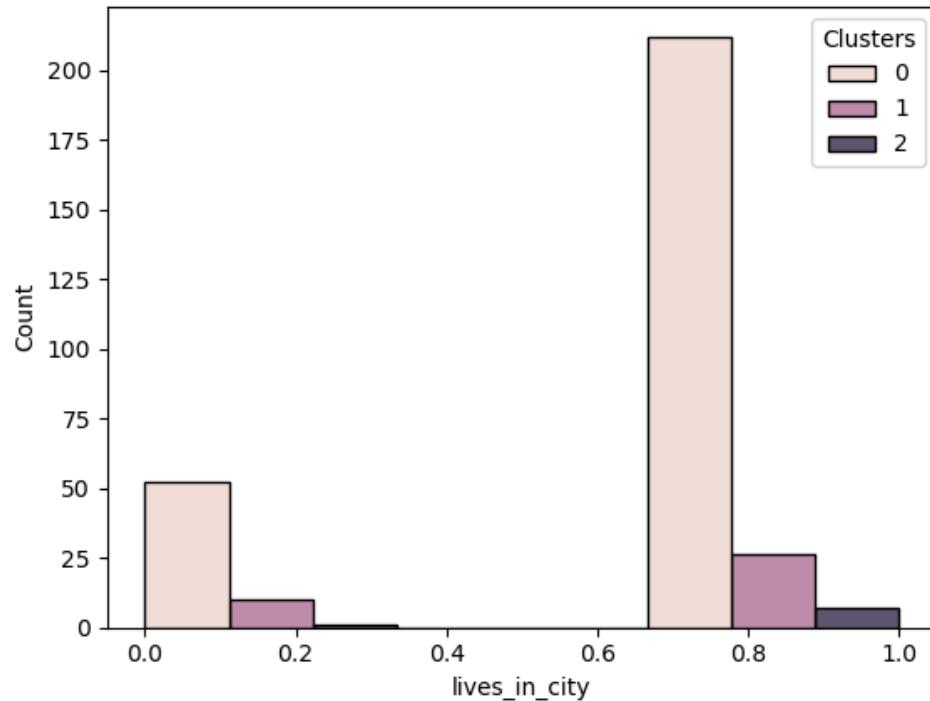
```
In [205]: sns.histplot(x = temp_student1['large_family'],
hue=temp_student1['Clusters'], multiple="dodge", bins=4)
```

Out[205]: <Axes: xlabel='large\_family', ylabel='Count'>



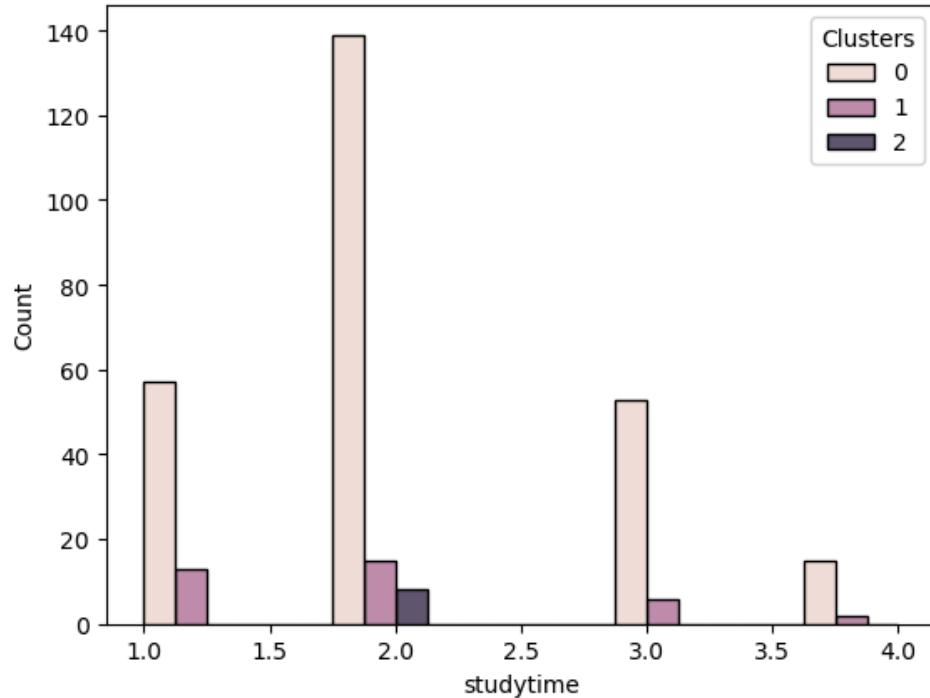
```
In [206]: sns.histplot(x = temp_student1['lives_in_city'],
hue=temp_student1['Clusters'], multiple="dodge", bins=3)
```

Out[206]: <Axes: xlabel='lives\_in\_city', ylabel='Count'>



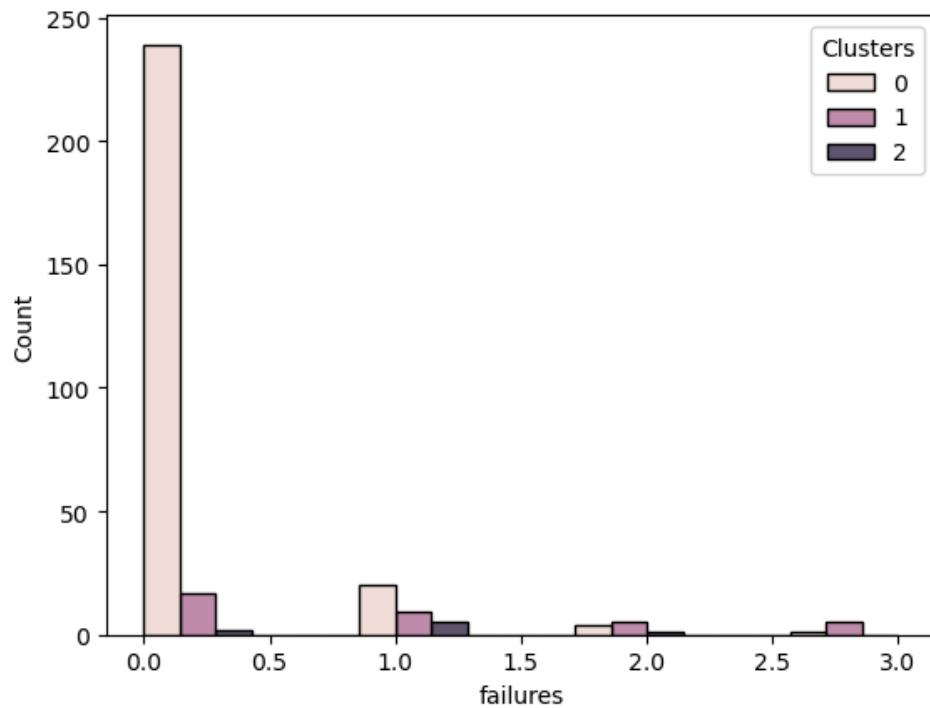
In [207]: `sns.histplot(x = temp_student1['studytime'], hue=temp_student1['Clusters'], multiple="dodge", bins=8)`

Out[207]: <Axes: xlabel='studytime', ylabel='Count'>



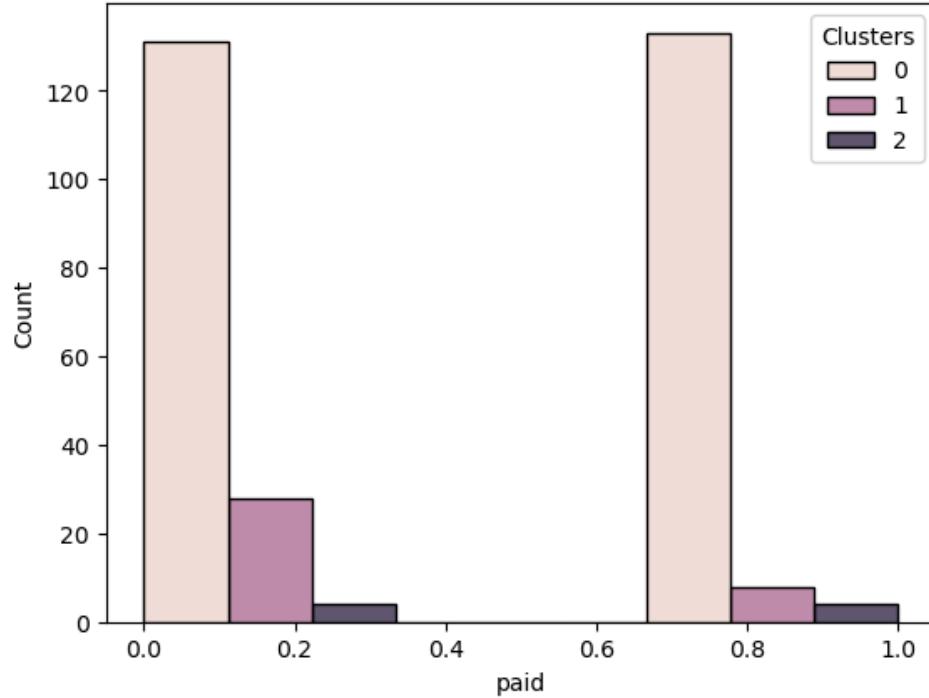
In [208]: `sns.histplot(x = temp_student1['failures'], hue=temp_student1['Clusters'], multiple="dodge", bins=7)`

Out[208]: <Axes: xlabel='failures', ylabel='Count'>



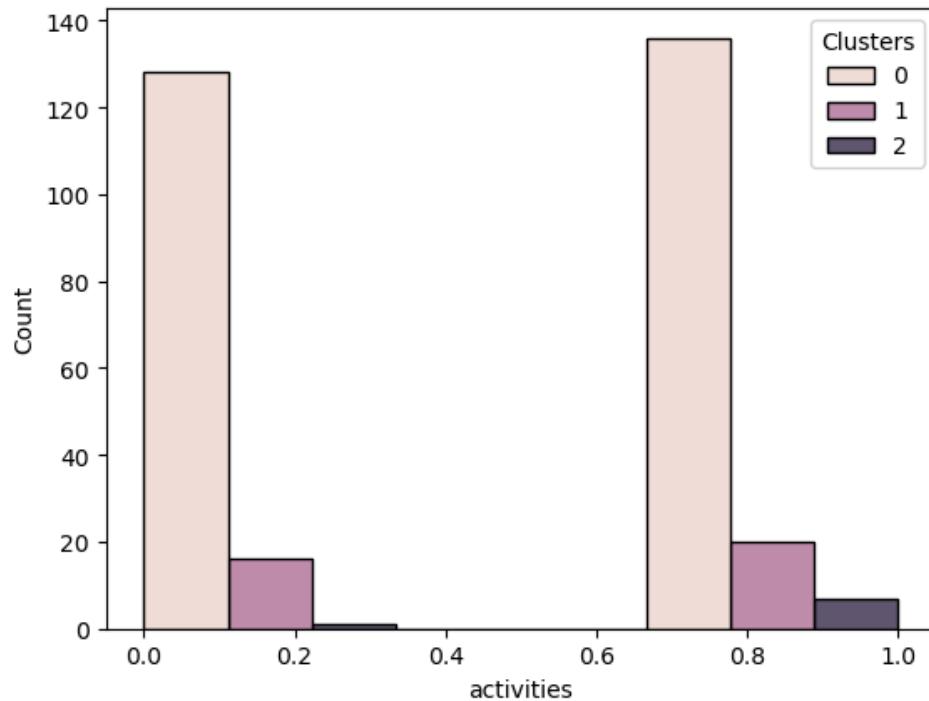
```
In [209]: sns.histplot(x = temp_student1['paid'], hue=temp_student1['Clusters'],  
multiple="dodge", bins=3)
```

Out[209]: <Axes: xlabel='paid', ylabel='Count'>



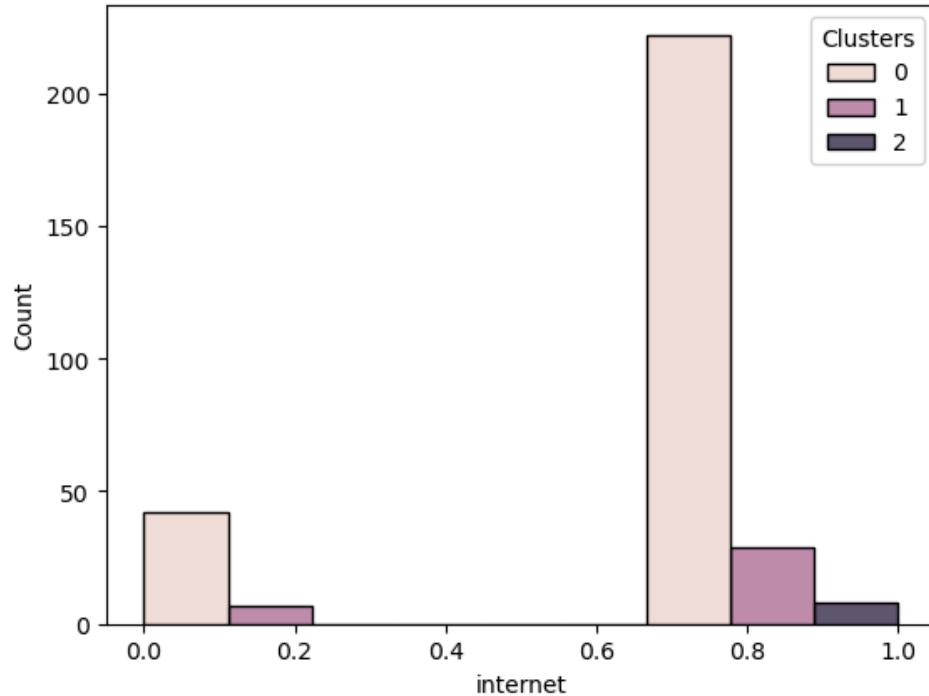
```
In [210]: sns.histplot(x = temp_student1['activities'],  
hue=temp_student1['Clusters'], multiple="dodge", bins=3)
```

Out[210]: <Axes: xlabel='activities', ylabel='Count'>



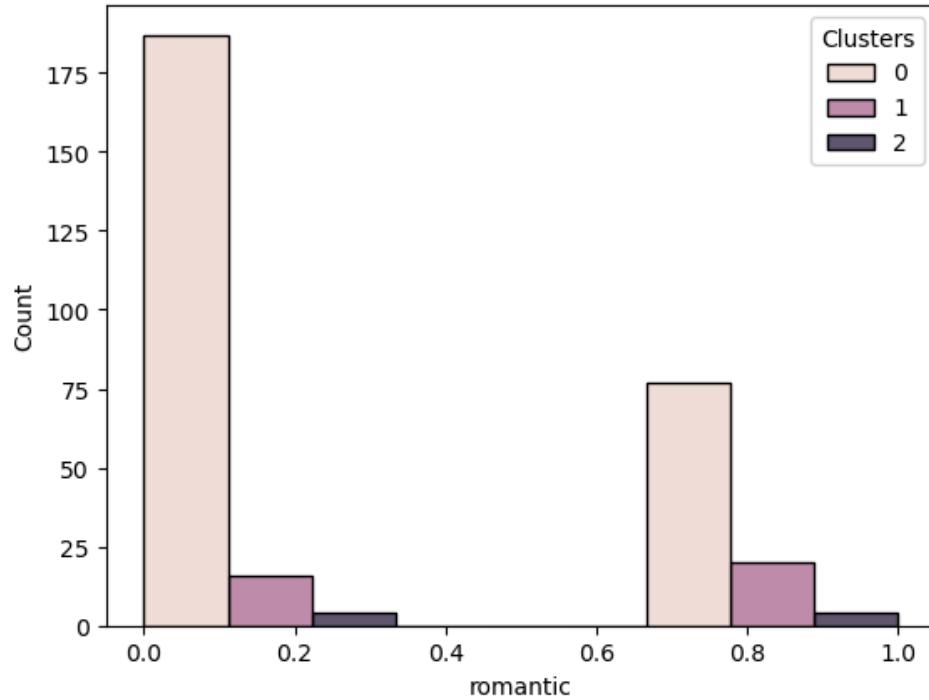
```
In [211]: sns.histplot(x = temp_student1['internet'],
hue=temp_student1['Clusters'], multiple="dodge", bins=3)
```

Out[211]: <Axes: xlabel='internet', ylabel='Count'>



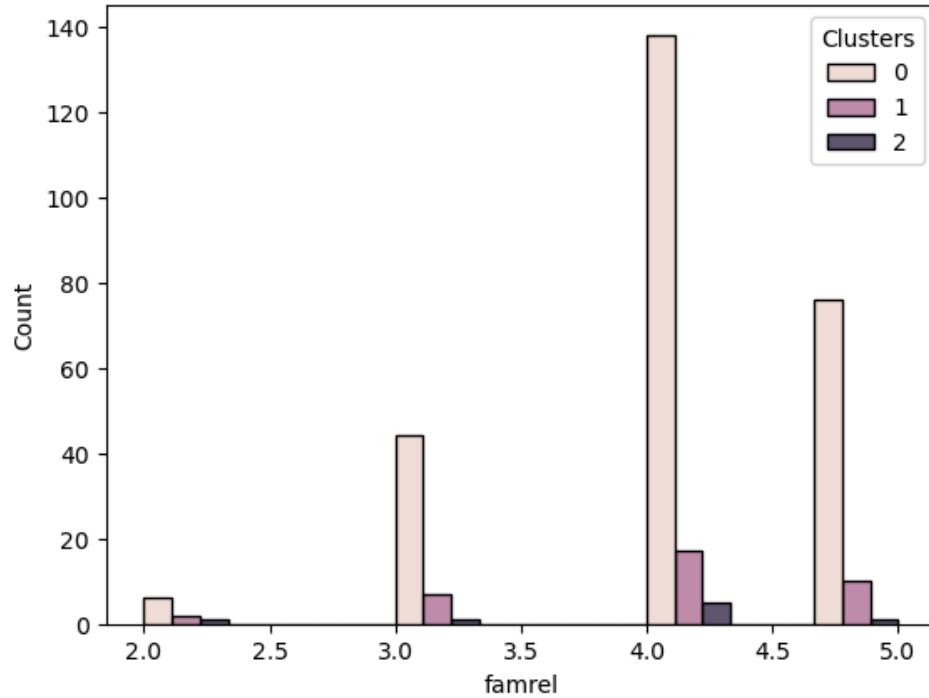
```
In [212]: sns.histplot(x = temp_student1['romantic'],
hue=temp_student1['Clusters'], multiple="dodge", bins=3)
```

Out[212]: <Axes: xlabel='romantic', ylabel='Count'>



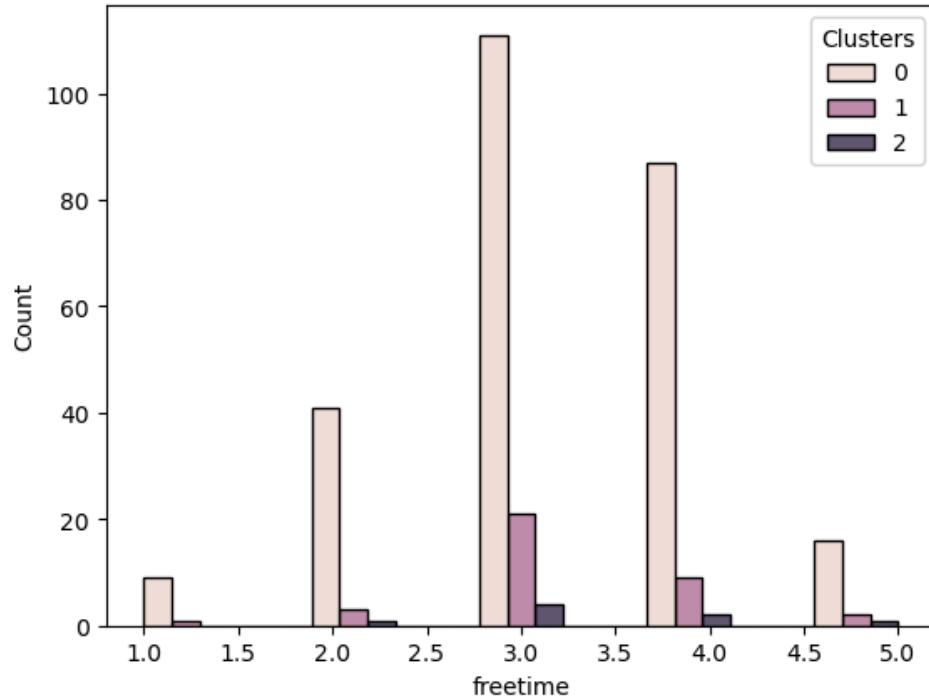
```
In [213]: sns.histplot(x = temp_student1['famrel'], hue=temp_student1['Clusters'],  
multiple="dodge", bins=9)
```

Out[213]: <Axes: xlabel='famrel', ylabel='Count'>



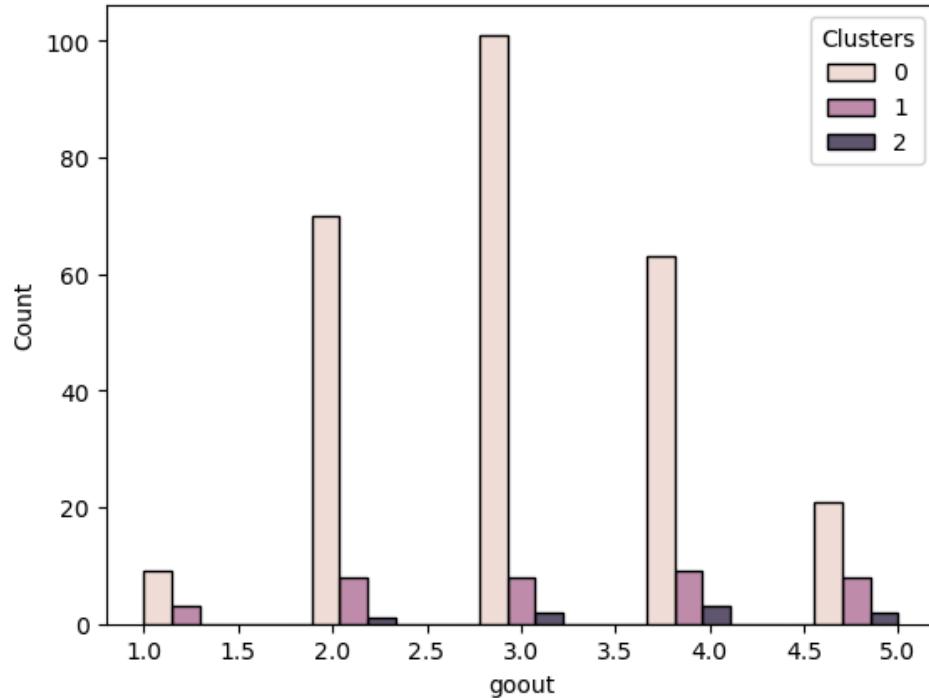
```
In [214]: sns.histplot(x = temp_student1['freetime'],  
hue=temp_student1['Clusters'], multiple="dodge", bins=9)
```

Out[214]: <Axes: xlabel='freetime', ylabel='Count'>



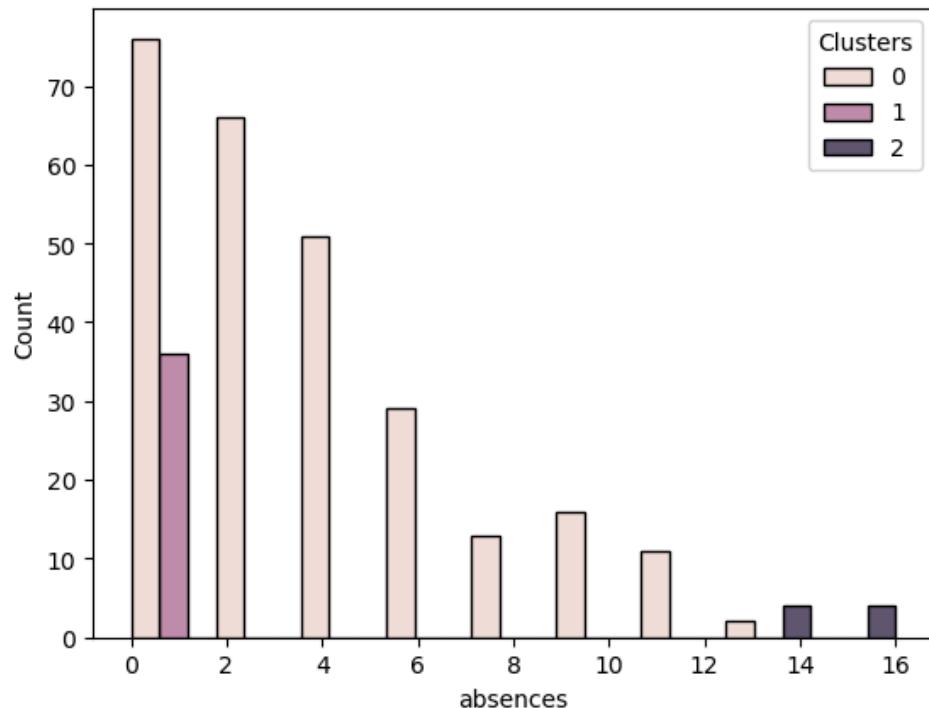
```
In [215]: sns.histplot(x = temp_student1['goout'], hue=temp_student1['Clusters'],  
multiple="dodge", bins=9)
```

Out[215]: <Axes: xlabel='goout', ylabel='Count'>



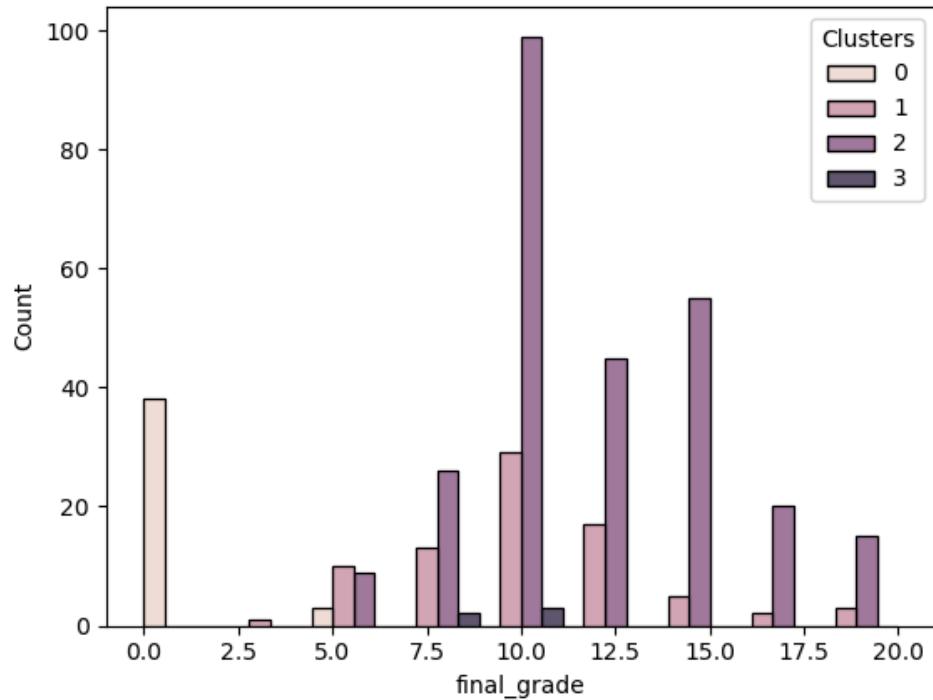
```
In [216]: sns.histplot(x = temp_student1['absences'],  
hue=temp_student1['Clusters'], multiple="dodge", bins=9)
```

Out[216]: <Axes: xlabel='absences', ylabel='Count'>



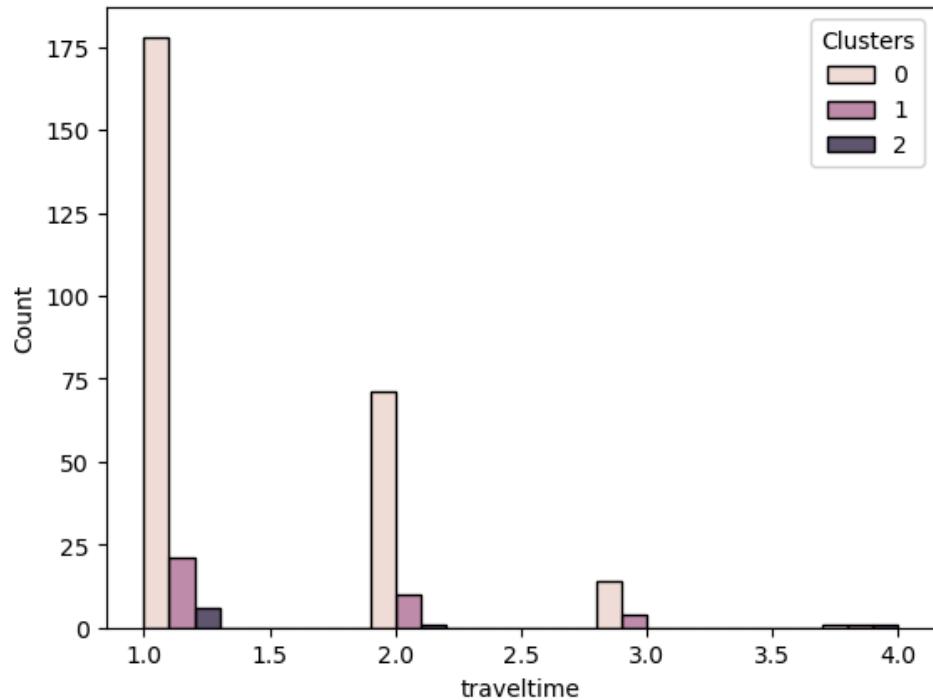
```
In [217]: sns.histplot(x = temp_student['final_grade'],  
hue=temp_student['Clusters'], multiple="dodge", bins=9)
```

Out[217]: <Axes: xlabel='final\_grade', ylabel='Count'>



```
In [218]: sns.histplot(x = temp_student1['traveltime'],  
hue=temp_student1['Clusters'], multiple="dodge", bins=10)
```

Out[218]: <Axes: xlabel='traveltime', ylabel='Count'>



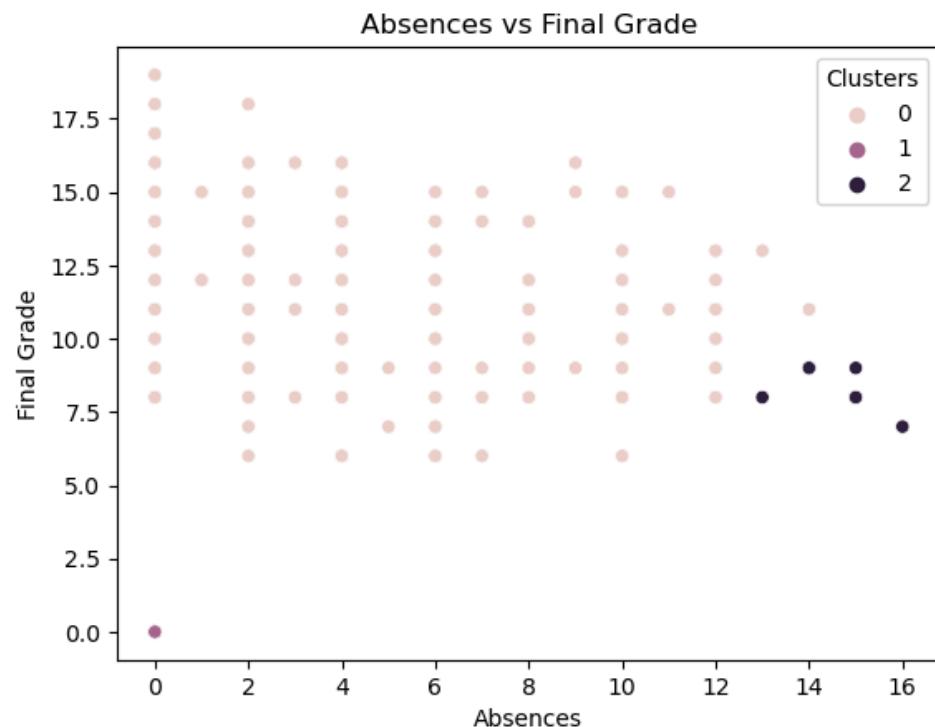
The Discriminatory variables are:

- final\_grade
- absences
- travelttime

## Scatter plots

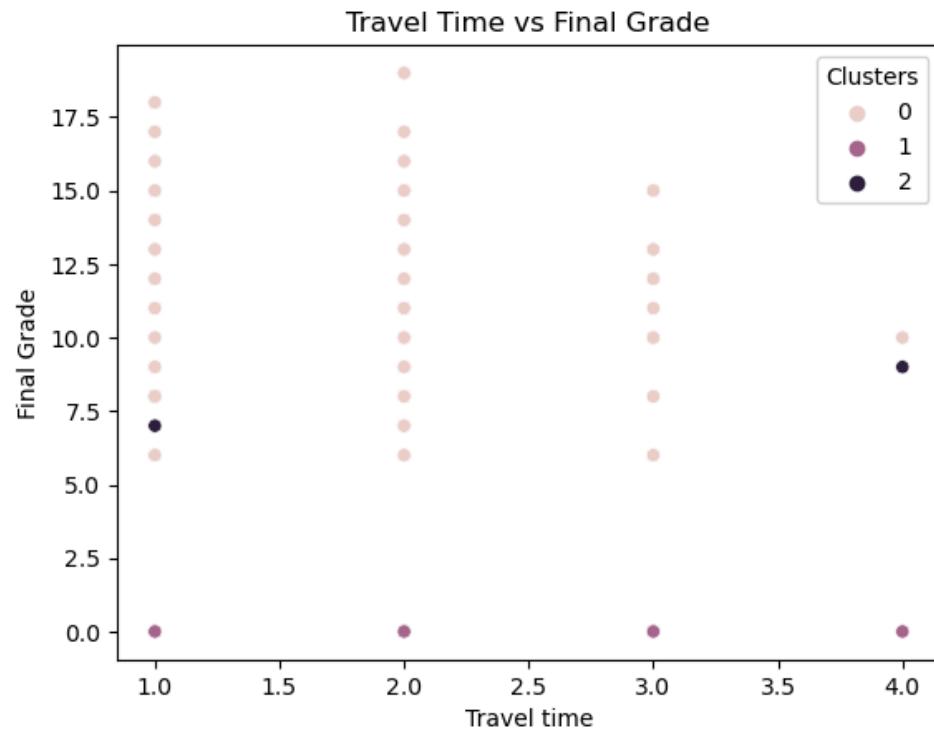
```
In #Scatter plot fpr absences and final grade
[219]: sns.scatterplot(x = temp_student1['absences'], y =
temp_student1['final_grade'], hue = temp_student1['Clusters'])
plt.xlabel('Absences')
plt.ylabel('Final Grade')
plt.title('Absences vs Final Grade')
```

Out[219]: Text(0.5, 1.0, 'Absences vs Final Grade')



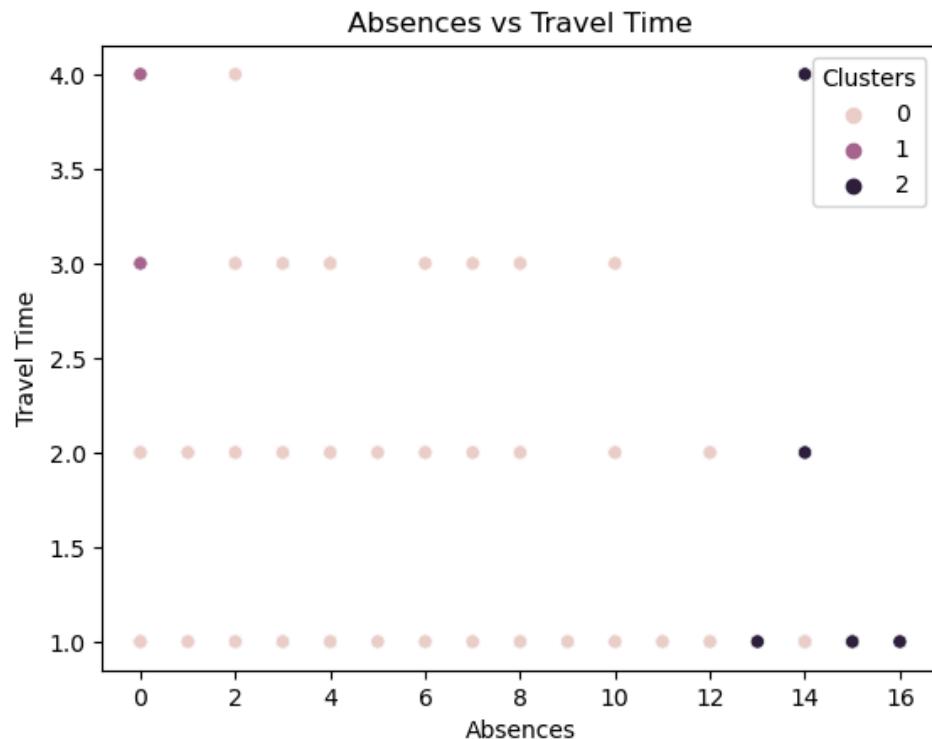
```
In [220]: #Scatter plot for travel time and final grade  
sns.scatterplot(x = temp_student1['traveltime'], y =  
temp_student1['final_grade'], hue = temp_student1['Clusters'])  
plt.xlabel('Travel time')  
plt.ylabel('Final Grade')  
plt.title('Travel Time vs Final Grade')
```

Out[220]: Text(0.5, 1.0, 'Travel Time vs Final Grade')



```
In [221]: #Scatter plot for absences and travel time  
sns.scatterplot(x = temp_student1['absences'], y =  
temp_student1['traveltime'], hue = temp_student1['Clusters'])  
plt.xlabel('Absences')  
plt.ylabel('Travel Time')  
plt.title('Absences vs Travel Time')
```

Out[221]: Text(0.5, 1.0, 'Absences vs Travel Time')



For DBSCAN, the cluster with most number of points is cluster 0. It does not look like there is any particular pattern for the clusters and for this dataset, this algorithm might not have been a good choice.

Looking at the plots for DBSCAN and K-means, it looks like K-Means was able to better categorise the students.

## Honour Code

I confirm that all work submitted is my own and that I have neither given, sought, nor received aid in relation to this assignment.